

torch

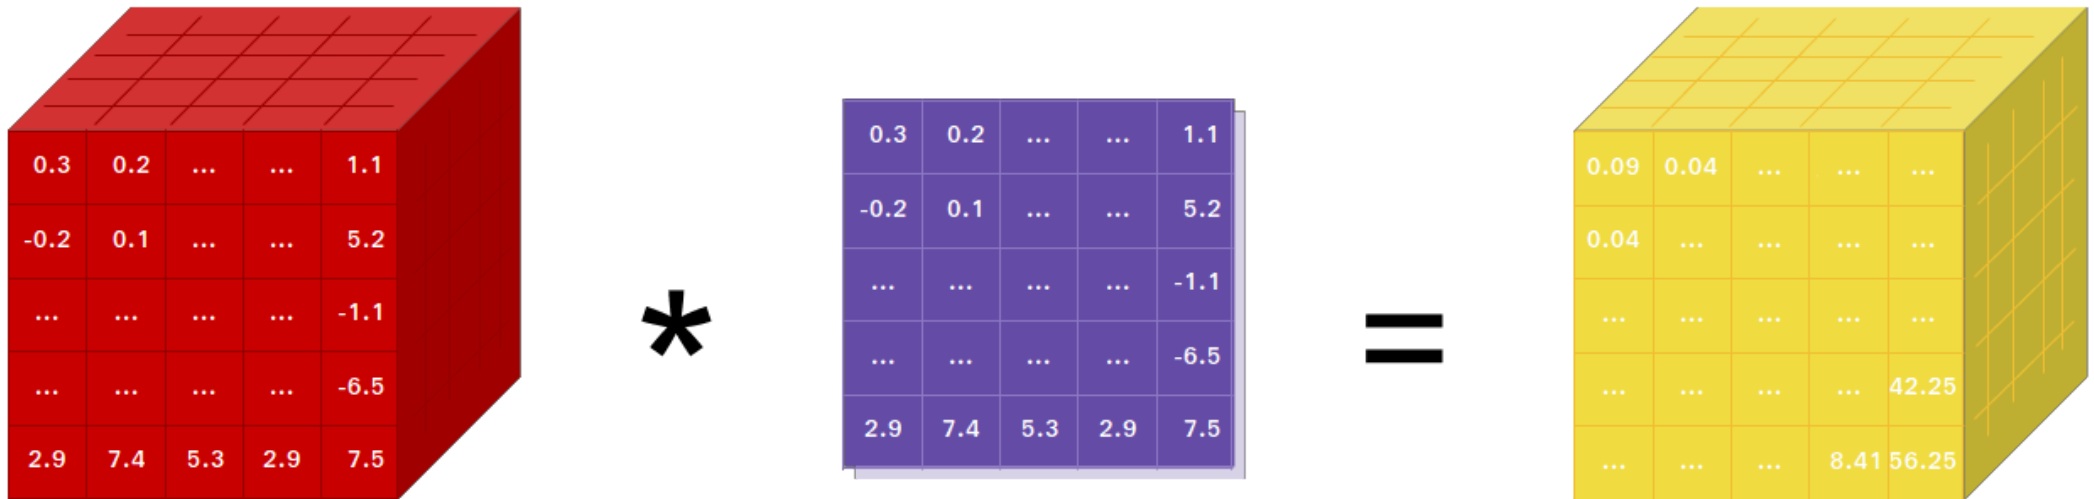
Lecture 17

Dr. Colin Rundel

PyTorch

PyTorch is a Python package that provides two high-level features:

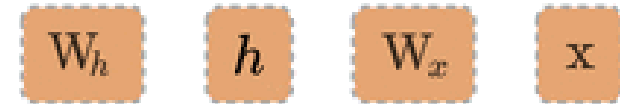
- Tensor computation (like NumPy) with strong GPU acceleration
- Deep neural networks built on a tape-based autograd system



```
1 import torch
2 torch.__version__
```

'2.6.0'

A graph is created on the fly



```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
```



Tensors

are the basic data abstraction in PyTorch and are implemented by the `torch.Tensor` class. They behave in much the same way as the other array libraries we've seen so far (`numpy`, `jax`, etc.)

```
1 torch.zeros(3)
```

```
tensor([0., 0., 0.])
```

```
1 torch.ones(3,2)
```

```
tensor([[1., 1.],
        [1., 1.],
        [1., 1.]])
```

```
1 torch.empty(2,2,2)
```

```
tensor([[[0., 0.],
         [0., 0.]],
        [[0., 0.],
         [0., 0.]])
```

```
1 torch.manual_seed(1234)
```

```
<torch._C.Generator object at
0x3015a3410>
```

```
1 torch.rand(2,2,2,2)
```

```
tensor([[[[0.02898, 0.40190],
          [0.25984, 0.36664]],
        [[0.05830, 0.70064],
          [0.05180, 0.46814]]],
        [[0.67381, 0.33146],
          [0.78371, 0.56306]],
        [[0.77485, 0.82080],
          [0.27928, 0.68171]]]])
```

Constants

As expected, tensors can be constructed from constant numeric values in lists or tuples.

```
1 torch.tensor(1)
```

```
tensor(1)
```

```
1 torch.tensor((1,2))
```

```
tensor([1, 2])
```

```
1 torch.tensor([[1,2,3], [4,5,6]])
```

```
tensor([[1, 2, 3],  
        [4, 5, 6]])
```

```
1 torch.tensor([(1,2,3), [4,5,6]])
```

```
tensor([[1, 2, 3],  
        [4, 5, 6]])
```

```
1 torch.tensor([(1,1,1), [4,5]])
```

```
ValueError: expected sequence of length 3  
at dim 1 (got 2)
```

```
1 torch.tensor([["A"]])
```

```
ValueError: too many dimensions 'str'
```

```
1 torch.tensor([[True]])
```

```
tensor([[True]])
```

Tensor Types

Data type	dtype	type()	Comment
32-bit float	float32 or float	FloatTensor	Default float
64-bit float	float64 or double	DoubleTensor	
16-bit float	float16 or half	HalfTensor	
16-bit brain float	bfloat16	BFloat16Tensor	
64-bit complex float	complex64		
128-bit complex float	complex128 or cdouble		
8-bit integer (unsigned)	uint8	ByteTensor	Default integer
8-bit integer (signed)	int8	CharTensor	
16-bit integer (signed)	int16 or short	ShortTensor	
32-bit integer (signed)	int32 or int	IntTensor	
64-bit integer (signed)	int64 or long	LongTensor	
Boolean	bool	BoolTensor	

Specifying types

Just like NumPy and Pandas, types are specified via the `dtype` argument and can be inspected via the `dtype` attribute.

```
1 a = torch.tensor([1,2,3]); a
```

```
tensor([1, 2, 3])
```

```
1 a.dtype
```

```
torch.int64
```

```
1 b = torch.tensor([1,2,3], dtype=torch.float16); b
```

```
tensor([1., 2., 3.], dtype=torch.float16)
```

```
1 b.dtype
```

```
torch.float16
```

```
1 c = torch.tensor([1.,2.,3.]); c
```

```
tensor([1., 2., 3.])
```

```
1 c.dtype
```

```
torch.float32
```

```
1 d = torch.tensor([1,2,3], dtype=torch.float64); d
```

```
tensor([1., 2., 3.], dtype=torch.float64)
```

```
1 d.dtype
```

```
torch.float64
```

Type precision

When using types with less precision it is important to be careful about underflow and overflow (ints) and rounding errors (floats).

```
1 torch.tensor([128], dtype=torch.int8)
```

RuntimeError: value cannot be converted to type int8 without overflow

```
1 torch.tensor([128]).to(torch.int8)
```

tensor([-128], dtype=torch.int8)

```
1 torch.tensor([255]).to(torch.uint8)
```

tensor([255], dtype=torch.uint8)

```
1 torch.tensor([300]).to(torch.uint8)
```

tensor([44], dtype=torch.uint8)

```
1 torch.tensor([300]).to(torch.int16)
```

tensor([300], dtype=torch.int16)

```
1 torch.tensor(1/3, dtype=torch.float16)
```

tensor(0.33325195, dtype=torch.float16)

```
1 torch.tensor(1/3, dtype=torch.float32)
```

tensor(0.33333334)

```
1 torch.tensor(1/3, dtype=torch.float64)
```

tensor(0.33333333, dtype=torch.float64)

```
1 torch.tensor(1/3, dtype=torch.bfloat16)
```

tensor(0.33398438, dtype=torch.bfloat16)

NumPy conversion

It is possible to easily move between NumPy arrays and Tensors via the `from_numpy()` function and `numpy()` method.

```
1 a = np.eye(3,3)
2 torch.from_numpy(a)
```

```
tensor([[1., 0., 0.],
        [0., 1., 0.],
        [0., 0., 1.]], dtype=torch.float64)
```

```
1 b = np.array([1,2,3])
2 torch.from_numpy(b)
```

```
tensor([1, 2, 3])
```

```
1 c = torch.rand(2,3)
2 c.numpy()
```

```
array([[0.28367, 0.65673, 0.23876],
       [0.73128, 0.60122, 0.30433]], dtype=float32)
```

```
1 d = torch.ones(2,2, dtype=torch.int64)
2 d.numpy()
```

```
array([[1, 1],
       [1, 1]])
```

Math & Logic

Just like NumPy torch `tensor` objects support basic mathematical and logical operations with scalars and other tensors - torch provides implementations of most commonly needed mathematical functions.

```
1 torch.ones(2,2) * 7 -1
```

```
tensor([[6., 6.],  
        [6., 6.]])
```

```
1 torch.ones(2,2) + torch.tensor([[1,2],
```

```
tensor([[2., 3.],  
        [4., 5.]])
```

```
1 2 ** torch.tensor([[1,2], [3,4]])
```

```
tensor([[ 2,  4],  
        [ 8, 16]])
```

```
1 2 ** torch.tensor([[1,2], [3,4]]) > 5
```

```
tensor([[False, False],  
        [ True,  True]])
```

```
1 x = torch.rand(2,2)
```

```
2 torch.ones(2,2) @ x
```

```
tensor([[1.22126317, 1.36931109],  
        [1.22126317, 1.36931109]])
```

```
1 torch.clamp(x*2-1, -0.5, 0.5)
```

```
tensor([[ -0.49049568,  0.25872374],  
        [ 0.50000000,  0.47989845]])
```

```
1 torch.mean(x)
```

```
tensor(0.64764357)
```

```
1 torch.sum(x)
```

```
tensor(2.59057426)
```

```
1 torch.min(x)
```

```
tensor(0.25475216)
```

Broadcasting

Like NumPy, cases where tensor dimensions do not match use the broadcasting heuristic.

The rules for broadcasting are:

- Each tensor must have at least one dimension - no empty tensors.
- Comparing the dimension sizes of the two tensors, going from last to first:
 - Each dimension must be equal, or
 - One of the dimensions must be of size 1, or
 - The dimension does not exist in one of the tensors

Exercise 1

Consider the following 6 tensors:

```
1 a = torch.rand(4, 3, 2)
2 b = torch.rand(3, 2)
3 c = torch.rand(2, 3)
4 d = torch.rand(0)
5 e = torch.rand(3, 1)
6 f = torch.rand(1, 2)
```

which of the above could be multiplied together and produce a valid result via broadcasting (e.g. $a*b$, $a*c$, $a*d$, etc.).

Explain why or why not broadcasting was able to be applied in each case.

```
1 countdown::countdown(5)
```

Inplace modification

In instances where we need to conserve memory it is possible to apply many functions such that a new tensor is not created but the original value(s) are replaced. These functions share the same name with the original functions but have a `_` suffix.

```
1 a = torch.rand(2,2)
2 print(a)
```

```
tensor([[0.31861043, 0.29080772],
        [0.41960979, 0.37281448]])
```

```
1 print(torch.exp(a))
```

```
tensor([[1.37521553, 1.33750737],
        [1.52136779, 1.45181489]])
```

```
1 print(a)
```

```
tensor([[0.31861043, 0.29080772],
        [0.41960979, 0.37281448]])
```

```
1 print(torch.exp_(a))
```

```
tensor([[1.37521553, 1.33750737],
        [1.52136779, 1.45181489]])
```

```
1 print(a)
```

```
tensor([[1.37521553, 1.33750737],
        [1.52136779, 1.45181489]])
```

Inplace arithmetic

All arithmetic functions are available as methods of the Tensor class,

```
1 a = torch.ones(2, 2)
2 b = torch.rand(2, 2)
```

```
1 a+b
```

```
tensor([[1.37689185, 1.01077938],
        [1.94549370, 1.76611161]])
```

```
1 print(a)
```

```
tensor([[1., 1.],
        [1., 1.]])
```

```
1 print(b)
```

```
tensor([[0.37689191, 0.01077944],
        [0.94549364, 0.76611167]])
```

```
1 a.add_(b)
```

```
tensor([[1.37689185, 1.01077938],
        [1.94549370, 1.76611161]])
```

```
1 print(a)
```

```
tensor([[1.37689185, 1.01077938],
        [1.94549370, 1.76611161]])
```

```
1 print(b)
```

```
tensor([[0.37689191, 0.01077944],
        [0.94549364, 0.76611167]])
```

Changing tensor shapes

The `shape` of a tensor can be changed using the `view()` or `reshape()` methods. The former guarantees that the result shares data with the original object (but requires contiguity), the latter may or may not copy the data.

```
1 x = torch.zeros(3, 2)
2 y = x.view(2, 3)
```

```
1 y
```

```
tensor([[0., 0., 0.],
        [0., 0., 0.]])
```

```
1 x.fill_(1)
```

```
tensor([[1., 1.],
        [1., 1.],
        [1., 1.]])
```

```
1 y
```

```
tensor([[1., 1., 1.],
        [1., 1., 1.]])
```

```
1 x = torch.zeros(3, 2)
2 y = x.t()
```

```
1 x.view(6)
```

```
tensor([0., 0., 0., 0., 0.,
        0.])
```

```
1 y.view(6)
```

RuntimeError: view size is not compatible with input tensor's size and stride (at least one dimension spans across two contiguous subspaces). Use `.reshape(...)` instead.

```
1 z = y.reshape(6)
2 x.fill_(1)
```

```
tensor([[1., 1.],
        [1., 1.],
        [1., 1.]])
```

```
1 y
```

```
tensor([[1., 1., 1.],
        [1., 1., 1.]])
```

```
1 z
```

```
tensor([0., 0., 0., 0., 0.,
        0.])
```

Adding or removing dimensions

The `squeeze()` and `unsqueeze()` methods can be used to remove or add length 1 dimension(s) to a tensor.

```
1 x = torch.zeros(1,3,1)
```

```
1 x.squeeze().shape
```

```
torch.Size([3])
```

```
1 x.squeeze(0).shape
```

```
torch.Size([3, 1])
```

```
1 x.squeeze(1).shape
```

```
torch.Size([1, 3, 1])
```

```
1 x.squeeze(2).shape
```

```
torch.Size([1, 3])
```

```
1 x = torch.zeros(3,2)
```

```
1 x.unsqueeze(0).shape
```

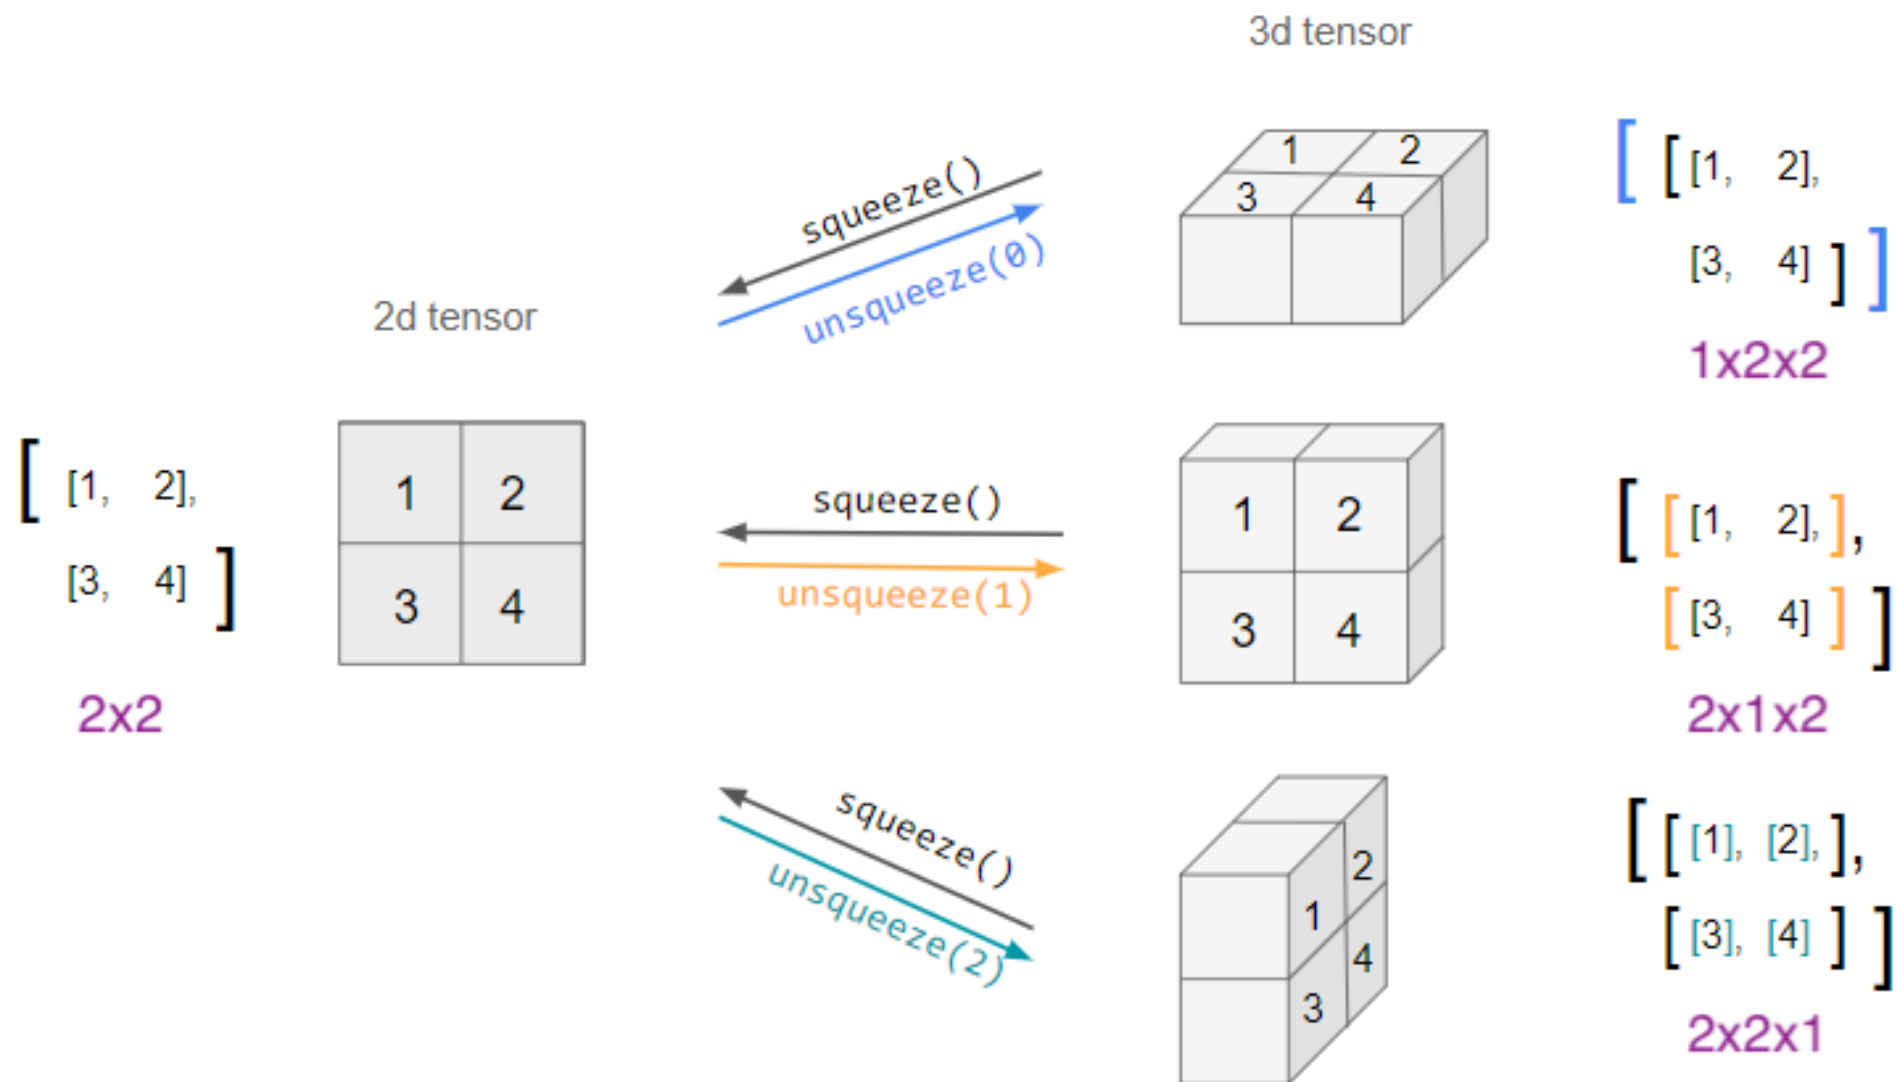
```
torch.Size([1, 3, 2])
```

```
1 x.unsqueeze(1).shape
```

```
torch.Size([3, 1, 2])
```

```
1 x.unsqueeze(2).shape
```

```
torch.Size([3, 2, 1])
```

Exercise 2

Given the following tensors,

```
1 a = torch.ones(4,3,2)
2 b = torch.rand(3)
3 c = torch.rand(5,3)
```

what reshaping is needed to make it possible so that $a * b$ and $a * c$ can be calculated via broadcasting?

```
1 countdown::countdown(3)
```

Autograd

Tensor expressions

Gradient tracking can be enabled using the `requires_grad` argument at initialization, alternatively the `requires_grad` flag can be set on the tensor or the `enable_grad()` context manager used (via `with`).

```
1 x = torch.linspace(0, 2, steps=21, requires_grad=True)
2 x
```

```
tensor([0.00000000, 0.10000000, 0.20000000, 0.30000001, 0.40000001, 0.50000000,
        0.60000002, 0.69999999, 0.80000001,
         0.90000004, 1.00000000, 1.10000002, 1.20000005, 1.29999995, 1.39999998,
        1.50000000,
         1.60000002, 1.70000005, 1.79999995, 1.89999998, 2.00000000],
        requires_grad=True)
```

```
1 y = 3*x + 2
2 y
```

```
tensor([2.00000000, 2.29999995, 2.59999990, 2.90000010, 3.20000005, 3.50000000,
        3.80000019, 4.09999990, 4.40000010,
         4.69999981, 5.00000000, 5.30000019, 5.60000038, 5.89999962, 6.19999981,
        6.50000000,
         6.80000019, 7.10000038, 7.39999962, 7.69999981, 8.00000000], grad_fn=
<AddBackward0>)
```

Computational graph

Basics of the computation graph can be explored via the `next_functions` attribute

```
1 y.grad_fn
```

```
<AddBackward0 object at 0x30c815390>
```

```
1 y.grad_fn.next_functions
```

```
((<MulBackward0 object at 0x30bb8e740>, 0), (None, 0))
```

```
1 y.grad_fn.next_functions[0][0].next_functions
```

```
((<AccumulateGrad object at 0x30c817f10>, 0), (None, 0))
```

```
1 y.grad_fn.next_functions[0][0].next_functions[0][0].next_functions
```

```
()
```

Autogradient

In order to calculate the gradients we use the `backward()` method on the *output* tensor (must be a scalar), this then makes the `grad` attribute available for the input (leaf) tensors.

```
1 out = y.sum()  
2 out.backward()  
3 out
```

```
tensor(105., grad_fn=<SumBackward0>)
```

```
1 y.grad
```

```
1 x.grad
```

```
tensor([3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,  
3., 3.])
```

A bit more complex

```
1 n = 21
2 x = torch.linspace(0, 2, steps=n, requires_grad=True)
3 m = torch.rand(n, requires_grad=True)
4
5 y = m*x + 2
6
7 y.backward(torch.ones(n))
```

```
1 x.grad
```

```
tensor([0.23227984, 0.72686875, 0.11874896, 0.39512146, 0.71987736, 0.75950843,
        0.53108865, 0.64494550, 0.72242016,
         0.44158769, 0.36338443, 0.88182861, 0.98741043, 0.73160070, 0.28143251,
        0.06507802,
         0.00649202, 0.50345892, 0.30815977, 0.37417805, 0.42968810])
```

```
1 m.grad
```

```
tensor([0.00000000, 0.10000000, 0.20000000, 0.30000001, 0.40000001, 0.50000000,
        0.60000002, 0.69999999, 0.80000001,
         0.90000004, 1.00000000, 1.10000002, 1.20000005, 1.29999995, 1.39999998,
        1.50000000,
         1.60000002, 1.70000005, 1.79999995, 1.89999998, 2.00000000])
```

High-level autograd API

allows for the automatic calculation and evaluation of the jacobian and hessian for a function defined using tensors.

```
1 def f(x, y):  
2     return 3*x + 1 + 2*y**2 + x*y
```

```
1 for x in [0.,1.]:  
2     for y in [0.,1.]:  
3         print("x =",x, "y = ",y)  
4         inputs = (torch.tensor([x]), torch.tensor([y]))  
5         print(torch.autograd.functional.jacobian(f, inputs),"\n")
```

```
x = 0.0 y = 0.0  
(tensor([[3.]]), tensor([[0.]])
```

```
x = 0.0 y = 1.0  
(tensor([[4.]]), tensor([[4.]])
```

```
x = 1.0 y = 0.0  
(tensor([[3.]]), tensor([[1.]])
```

```
x = 1.0 y = 1.0  
(tensor([[4.]]), tensor([[5.]])
```



```
1 inputs = (torch.tensor([0.]), torch.tensor([0.]))
2 torch.autograd.functional.hessian(f, inputs)
```

```
((tensor([[0.]]), tensor([[1.]])), (tensor([[1.]]), tensor([[4.]])))
```

```
1 inputs = (torch.tensor([1.]), torch.tensor([1.]))
2 torch.autograd.functional.hessian(f, inputs)
```

```
((tensor([[0.]]), tensor([[1.]])), (tensor([[1.]]), tensor([[4.]])))
```

Demo 1 - Linear Regression w/ PyTorch

A basic model

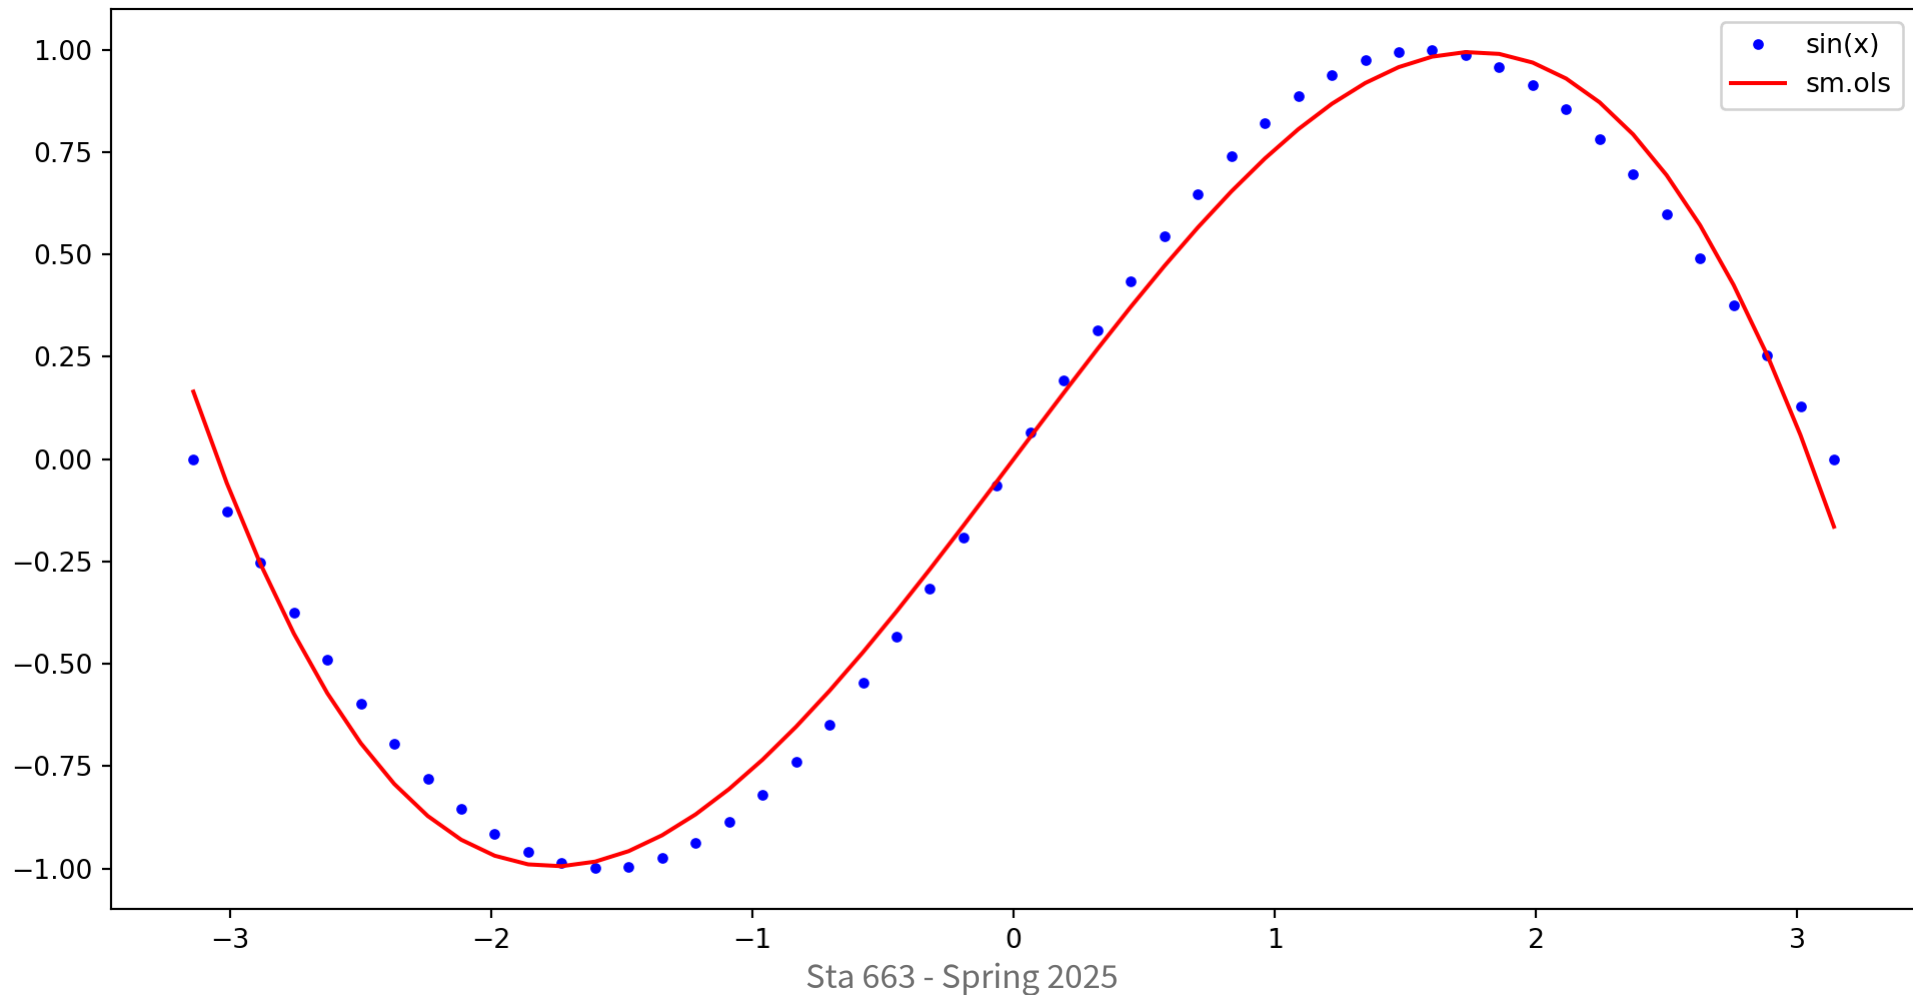
```
1 x = np.linspace(-math.pi, math.pi, 50)
2 y = np.sin(x)
3
4 lm = smf.ols(
5     "y~x+I(x**2)+I(x**3)",
6     data=pd.DataFrame({"x": x, "y": y})
7 ).fit()
8
9 print(lm.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	y	R-squared:	0.990			
Model:	OLS	Adj. R-squared:	0.989			
Method:	Least Squares	F-statistic:	1455.			
Date:	Wed, 19 Mar 2025	Prob (F-statistic):	1.44e-45			
Time:	09:33:38	Log-Likelihood:	60.967			
No. Observations:	50	AIC:	-113.9			
Df Residuals:	46	BIC:	-106.3			
Df Model:	3					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	-7.958e-17	0.016	-5.03e-15	1.000	-0.032	0.032
x	0.8476	0.014	59.444	0.000	0.819	0.876
I(x ** 2)	3.692e-17	0.003	1.07e-14	1.000	-0.007	0.007
I(x ** 3)	-0.0912	0.002	-42.977	0.000	-0.095	-0.087
=====						

Predictions

```
1 plt.figure(figsize=(10,5), layout="constrained")
2 plt.plot(x, y, ".b", label="sin(x)")
3 plt.plot(x, lm.predict(), "-r", label="sm.ols")
4 plt.legend()
5 plt.show()
```



Making tensors

```
1 yt = torch.tensor(y)
2 Xt = torch.tensor(lm.model.exog)
3 bt = torch.randn((Xt.shape[1], 1), dtype=torch.float64, requires_grad=True)
```

```
1 yt.shape
```

```
torch.Size([50])
```

```
1 Xt.shape
```

```
torch.Size([50, 4])
```

```
1 bt.shape
```

```
torch.Size([4, 1])
```

```
1 yt_pred = (Xt @ bt).squeeze()
```

```
1 loss = (yt_pred - yt).pow(2).sum()
2 loss.item()
```

```
2119.277704016523
```

Gradient descent

```
1 learning_rate = 1e-6
2
3 loss.backward() # Compute the backward pass
4
5 with torch.no_grad():
6     bt -= learning_rate * bt.grad # Make the step
7
8     bt.grad = None # Reset the gradients
```

```
1 yt_pred = (Xt @ bt).squeeze()
2 loss = (yt_pred - yt).pow(2).sum()
3 loss.item()
```

2069.4881821807053

Putting it together

```
1 yt = torch.tensor(y).unsqueeze(1)
2 Xt = torch.tensor(lm.model.exog)
3 bt = torch.randn((Xt.shape[1], 1), dtype=torch.float64, requires_grad=True)
4
5 learning_rate = 1e-5
6 for i in range(5001):
7
8     yt_pred = Xt @ bt
9
10    loss = (yt_pred - yt).pow(2).sum()
11    if i % 500 == 0:
12        print(f"Step: {i},\tloss: {loss.item()}")
13
14    loss.backward()
15
16    with torch.no_grad():
17        bt -= learning_rate * bt.grad
18        bt.grad = None
```

Putting it together

```
Step: 0,      loss: 70161.1580804254
Step: 500,    loss: 14.791178300540242
Step: 1000,   loss: 8.825181658035252
Step: 1500,   loss: 5.311942717260374
Step: 2000,   loss: 3.2416251317783518
Step: 2500,   loss: 2.020671792951764
Step: 3000,   loss: 1.300022038356929
Step: 3500,   loss: 0.8742816442183533
Step: 4000,   loss: 0.6225166364100523
Step: 4500,   loss: 0.473473387453477
Step: 5000,   loss: 0.38513809895450724
```

```
1 print(bt)
```

```
tensor([[ 0.03141952],
        [ 0.78487683],
        [-0.00520719],
        [-0.08261045]], dtype=torch.float64,
requires_grad=True)
```

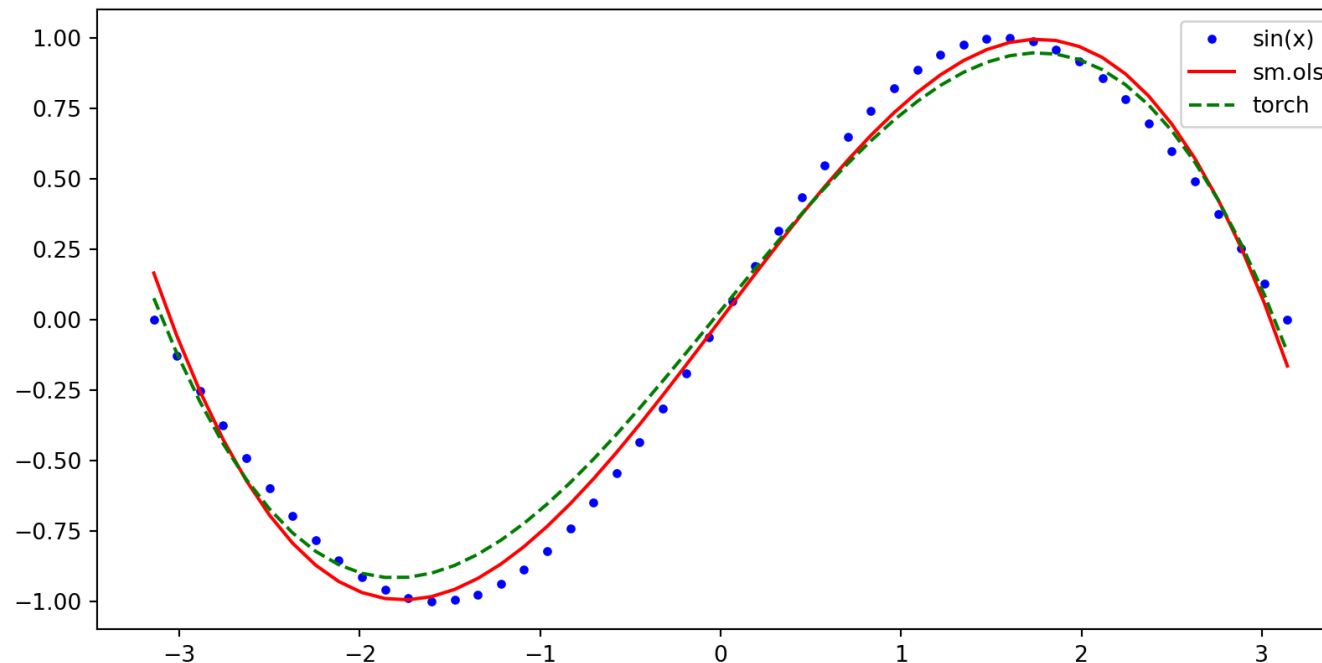

Comparing results

1 lm.params

```
Intercept    -7.958044e-17  
x             8.476289e-01  
I(x ** 2)     3.691708e-17  
I(x ** 3)    -9.120167e-02  
dtype: float64
```

1 bt

```
tensor([[ 0.03141952],  
        [ 0.78487683],  
        [-0.00520719],  
        [-0.08261045]],  
        dtype=torch.float64, requires_grad=True)
```



Demo 2 - Using a torch model

A simple model

```
1 class Model(torch.nn.Module):
2     def __init__(self, beta):
3         super().__init__()
4         beta.requires_grad = True
5         self.beta = torch.nn.Parameter(beta)
6
7     def forward(self, X):
8         return X @ self.beta
9
10 def training_loop(model, X, y, optimizer, n=1000):
11     losses = []
12     for i in range(n):
13         y_pred = model(X)
14
15         loss = (y_pred.squeeze() - y.squeeze()).pow(2).sum()
16         loss.backward()
17
18         optimizer.step()
19         optimizer.zero_grad()
20
21         losses.append(loss.item())
22
23     return losses
```

Fitting

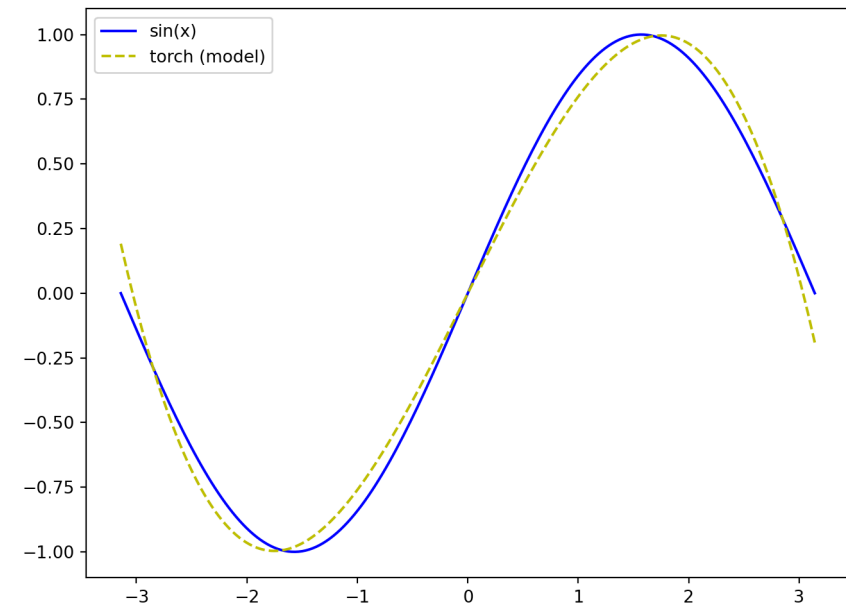
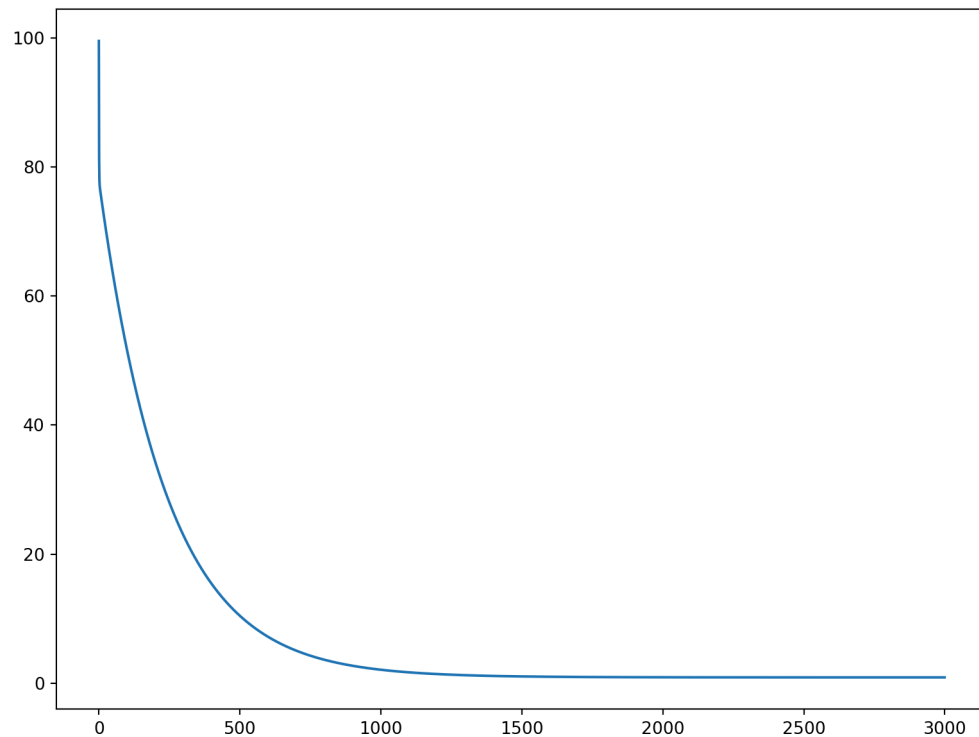
```
1 x = torch.linspace(-math.pi, math.pi, 200)
2 y = torch.sin(x)
3
4 X = torch.vstack((
5     torch.ones_like(x),
6     x,
7     x**2,
8     x**3
9 )) .T
10
11 m = Model(beta = torch.zeros(4))
12 opt = torch.optim.SGD(m.parameters(), lr=1e-5)
13
14 losses = training_loop(m, X, y, opt, n=3000)
```

Results

1 m.beta

Parameter containing:

```
tensor([-4.07514189e-10,  8.52953434e-01,  1.22972355e-10, -9.25917700e-02],  
requires_grad=True)
```

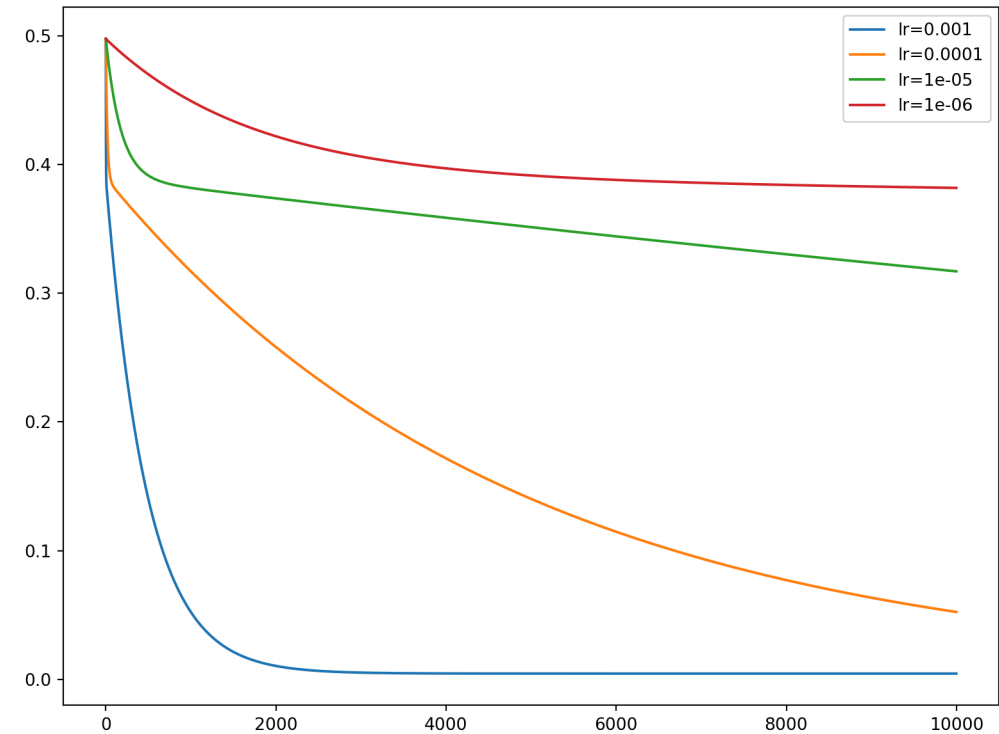


An all-in-one model

```
1 class Model(torch.nn.Module):
2     def __init__(self, X, y, beta=None):
3         super().__init__()
4         self.X = X
5         self.y = y
6         if beta is None:
7             beta = torch.zeros(X.shape[1])
8             beta.requires_grad = True
9             self.beta = torch.nn.Parameter(beta)
10
11     def forward(self, X):
12         return X @ self.beta
13
14     def fit(self, opt, n=1000, loss_fn = torch.nn.MSELoss()):
15         losses = []
16         for i in range(n):
17             loss = loss_fn(self.forward(self.X).squeeze(), self.y.squeeze())
18             loss.backward()
19             opt.step()
20             opt.zero_grad()
21             losses.append(loss.item())
```

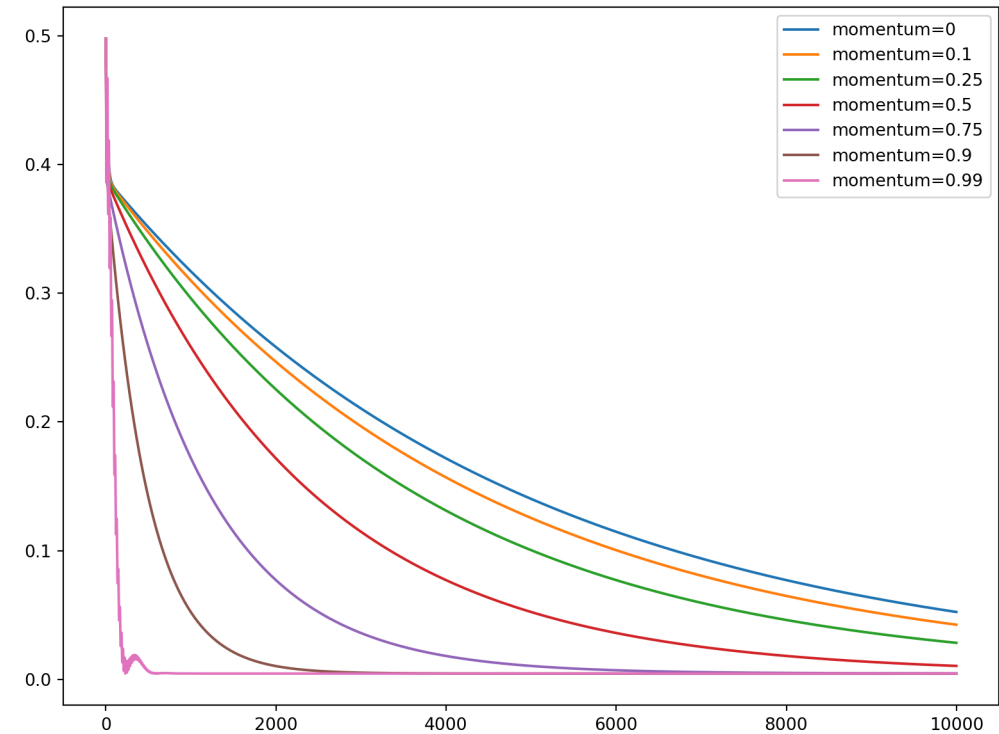
Learning rate and convergence

```
1 plt.figure(figsize=(8,6), layout="constrained")
2
3 for lr in [1e-3, 1e-4, 1e-5, 1e-6]:
4     m = Model(X, y)
5     opt = torch.optim.SGD(m.parameters())
6     losses = m.fit(opt, n=10000)
7
8     plt.plot(losses, label=f"{lr}")
9
10 plt.legend()
11 plt.show()
```



Momentum and convergence

```
1 plt.figure(figsize=(8,6), layout="constrained")
2
3 for momentum in [0, 0.1, 0.25, 0.5, 0.75, 0.9, 0.99]:
4     m = Model(X, y)
5     opt = torch.optim.SGD(
6         m.parameters(),
7         lr = 1e-4,
8         momentum = momentum
9     )
10    losses = m.fit(opt, n=10000)
11
12    plt.plot(losses, label=f"{momentum}")
13
14 plt.legend()
15 plt.show()
```



Optimizers and convergence

```
1 plt.figure(figsize=(8,6), layout="constrained")
2
3 opts = (torch.optim.SGD,
4         torch.optim.Adam,
5         torch.optim.Adagrad)
6
7 for opt_fn in opts:
8     m = Model(X, y)
9     opt = opt_fn(m.parameters(), lr=1e-4)
10    losses = m.fit(opt, n=10000)
11
12    plt.plot(losses, label=f"{opt_fn}")
13
14 plt.legend()
15 plt.show()
```

