

Optimization (cont.)

Lecture 14

Dr. Colin Rundel

Method Variants

Method: CG in scipy

Scipy's optimize module implements the conjugate gradient algorithm by Polak and Ribiere, a variant that does not require the Hessian,

- α_k is calculated via a line search along the direction p_k
- We replace the previous definition

$$\beta_{k+1} = \frac{r_{k+1}^T \nabla^2 f(x_k) p_k}{p_k^T \nabla^2 f(x_k) p_k}$$

with

$$\beta_{k+1}^{PR} = \frac{\nabla f(x_{k+1}) (\nabla f(x_{k+1}) - \nabla f(x_k))}{\nabla f(x_k)^T \nabla f(x_k)}$$

Method: Newton-CG & BFGS

These are both variants of Newton's method but do not require the Hessian (but can be used by the former if provided).

- The Broyden-Fletcher-Goldfarb-Shanno (BFGS) algorithm is a quasi-newton which iteratively approximates Hessian

Method: Nelder-Mead

This is a gradient free method that uses a series of simplexes which are used to iteratively bracket the minimum.

Method Summary

SciPy Method	Description	Gradient	Hessian
—	Gradient Descent (naive w/ backtracking)	✓	✗
—	Newton's method (naive w/ backtracking)	✓	✓
—	Conjugate Gradient (naive)	✓	✓
"CG"	Nonlinear Conjugate Gradient (Polak and Ribiere variation)	✓	✗
"Newton-CG"	Truncated Newton method (Newton w/ CG step direction)	✓	Optional
"BFGS"	Broyden, Fletcher, Goldfarb, and Shanno (Quasi-newton method)	Optional	✗
"L-BFGS-B"	Limited-memory BFGS (Quasi-newton method)	Optional	✗
"Nelder-Mead"	Nelder-Mead simplex reflection method	✗	✗

Methods collection


```
1 def define_methods(x, f, grad, hess, tol=1e-8):
2     return {
3         "naive_newton": lambda: newtons_method(x, f, grad, hess, tol=tol),
4         "naive_cg": lambda: conjugate_gradient(x, f, grad, hess, tol=tol),
5         "CG": lambda: optimize.minimize(f, x, jac=grad, method="CG", tol=tol),
6         "newton-cg": lambda: optimize.minimize(f, x, jac=grad, hess=None, method="newton-cg", tol=tol),
7         "newton-cg w/ H": lambda: optimize.minimize(f, x, jac=grad, hess=hess, method="newton-cg w/ H", tol=tol),
8         "bfgs": lambda: optimize.minimize(f, x, jac=grad, method="BFGS", tol=tol),
9         "bfgs w/o G": lambda: optimize.minimize(f, x, method="BFGS", tol=tol),
10        "l-bfgs-b": lambda: optimize.minimize(f, x, method="L-BFGS-B", tol=tol),
11        "nelder-mead": lambda: optimize.minimize(f, x, method="Nelder-Mead", tol=tol),
12    }
```

Method Timings

```
1 x = (1.6, 1.1)
2 f, grad, hess = mk_quad(0.7)
3 methods = define_methods(x, f, grad, hess)
4
5 df = pd.DataFrame({
6     key: timeit.Timer(methods[key]).repeat(10, 100) for key in methods
7 })
8
9 df
```

	naive_newton	naive_cg	CG	...	bfgs w/o G	l-bfgs-b	nelder-mead
0	0.029418	0.031432	0.010260	...	0.048270	0.029618	0.100661
1	0.028514	0.032393	0.009922	...	0.048328	0.026831	0.099941
2	0.028562	0.032846	0.009725	...	0.048308	0.026727	0.100675
3	0.029202	0.032913	0.009430	...	0.048968	0.026292	0.099893
4	0.029238	0.031702	0.009614	...	0.048761	0.026301	0.100530
5	0.029393	0.032050	0.009654	...	0.048400	0.026221	0.100207
6	0.029281	0.032417	0.009676	...	0.048008	0.026354	0.100014
7	0.029246	0.031986	0.009658	...	0.048347	0.026398	0.098695
8	0.029418	0.032260	0.009675	...	0.051405	0.026557	0.102137
9	0.029312	0.032251	0.009771	...	0.048857	0.026232	0.098924

[10 rows x 9 columns]



Timings across cost functions


```

1 def time_cost_func(n, x, name, cost_func, *args):
2     x = (1.6, 1.1)
3     f, grad, hess = cost_func(*args)
4     methods = define_methods(x, f, grad, hess)
5
6     return ( pd.DataFrame({
7         key: timeit.Timer(
8             methods[key]
9         ).repeat(n, n)
10        for key in methods
11    })
12    .melt()
13    .assign(cost_func = name)
14 )
15
16 df = pd.concat([
17     time_cost_func(10, x, "Well-cond quad", ml),
18     time_cost_func(10, x, "Ill-cond quad", mk),
19     time_cost_func(10, x, "Rosenbrock", mk_ro)
20 ])
21
22 df

```


	variable	value	cost_func
0	naive_newton	0.003061	Well-cond quad
1	naive_newton	0.003014	Well-cond quad
2	naive_newton	0.002951	Well-cond quad
3	naive_newton	0.002920	Well-cond quad
4	naive_newton	0.002978	Well-cond quad
..
85	nelder-mead	0.016167	Rosenbrock
86	nelder-mead	0.016182	Rosenbrock
87	nelder-mead	0.016139	Rosenbrock
88	nelder-mead	0.016248	Rosenbrock
89	nelder-mead	0.016163	Rosenbrock

[270 rows x 3 columns]



Random starting locations

```
1 x0s = np.random.default_rng(seed=1234).uniform(-2,2, (20,2))
2
3 df = pd.concat([
4     pd.concat([
5         time_cost_func(3, x0, "Well-cond quad", mk_quad, 0.7),
6         time_cost_func(3, x0, "Ill-cond quad", mk_quad, 0.02),
7         time_cost_func(3, x0, "Rosenbrock", mk_rosenbrock)
8     ])
9     for x0 in x0s
10 ])
```



Profiling - BFGS (cProfile)

```
1 import cProfile
2
3 f, grad, hess = mk_quad(0.7)
4 def run():
5     optimize.minimize(fun = f, x0 = (1.6, 1.1), jac=grad, method="BFGS", tol=1e-11)
6
7 cProfile.run('run()', sort="tottime")
```

1011 function calls in 0.001 seconds

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.001	0.001	_optimize.py:1328(_minimize_bfgs)
58	0.000	0.000	0.000	0.000	{method 'reduce' of 'numpy.ufunc' objects}
10	0.000	0.000	0.000	0.000	<string>:3(f)
26	0.000	0.000	0.000	0.000	_optimize.py:194(vecnorm)
9	0.000	0.000	0.000	0.000	_dcsrc.py:201(__call__)
1	0.000	0.000	0.001	0.001	_minimize.py:53(minimize)
36	0.000	0.000	0.000	0.000	fromnumeric.py:69(_wrapreduction)
20	0.000	0.000	0.000	0.000	numeric.py:2475(array_equal)
10	0.000	0.000	0.000	0.000	<string>:9(gradient)
9	0.000	0.000	0.000	0.000	_linesearch.py:100(scalar_search_wolfe1)
18	0.000	0.000	0.000	0.000	_dcsrc.py:269(_iterate)
9	0.000	0.000	0.000	0.000	_linesearch.py:86(derphi)
26	0.000	0.000	0.000	0.000	fromnumeric.py:2228(cumsum)

Profiling - BFGS (pyinstrument)

```
1 from pyinstrument import Profiler
2
3 f, grad, hess = mk_quad(0.7)
4
5 profiler = Profiler(interval=0.00001)
6
7 profiler.start()
8 opt = optimize.minimize(fun = f, x0 = (1.6, 1.1), jac=grad, method="BFGS", tol=1e-11)
9 p = profiler.stop()
10
11 profiler.print(show_all=True)
```

Profiling - Nelder-Mead

```
1 from pyinstrument import Profiler
2
3 f, grad, hess = mk_quad(0.7)
4
5 profiler = Profiler(interval=0.00001)
6
7 profiler.start()
8 opt = optimize.minimize(fun = f, x0 = (1.6, 1.1), method="Nelder-Mead", tol=1e-11)
9 p = profiler.stop()
10
11 profiler.print(show_all=True)
```

optimize.minimize() output

```
1 f, grad, hess = mk_quad(0.7)
```

```
1 optimize.minimize(fun = f, x0 = (1.6,  
2 jac=grad, method="BF")
```

```
message: Optimization terminated  
successfully.  
success: True  
status: 0  
fun: 1.2739256453439323e-11  
x: [-5.318e-07 -8.843e-06]  
nit: 6  
jac: [-3.510e-07 -2.860e-06]  
hess_inv: [[ 1.515e+00 -3.438e-03]  
            [-3.438e-03  3.035e+00]]  
nfev: 7  
njev: 7
```

```
1 optimize.minimize(fun = f, x0 = (1.6,  
2 jac=grad, hess=hess,  
3 method="Newton-CG")
```

```
message: Optimization terminated  
successfully.  
success: True  
status: 0  
fun: 2.3418652989289317e-12  
x: [ 0.000e+00  3.806e-06]  
nit: 11  
jac: [ 0.000e+00  4.102e-06]  
nfev: 12  
njev: 12  
nhev: 11
```


optimize.minimize() output (cont.)

```
1 optimize.minimize(fun = f, x0 = (1.6,  
2 jac=grad, method="CG"
```

```
message: Optimization terminated  
successfully.  
success: True  
status: 0  
fun: 1.4450021261144105e-32  
x: [-1.943e-16 -1.110e-16]  
nit: 2  
jac: [-1.282e-16 -3.590e-17]  
nfev: 5  
njev: 5
```

```
1 optimize.minimize(fun = f, x0 = (1.6,  
2 jac=grad, method="Ne
```

```
message: Optimization terminated  
successfully.  
success: True  
status: 0  
fun: 2.3077013477040082e-10  
x: [ 1.088e-05  3.443e-05]  
nit: 46  
nfev: 89  
final_simplex: (array([[ 1.088e-05,  
3.443e-05],  
[ 1.882e-05,  
-3.825e-05],  
[-3.966e-05,  
-3.147e-05]]), array([ 2.308e-10,  
3.534e-10,  6.791e-10]))
```

Collect

```

1 def run_collect(name, x0, cost_func, *
2   f, grad, hess = cost_func(*args)
3   methods = define_methods(x0, f, grad
4
5   res = []
6   for method in methods:
7       if method in skip: continue
8
9       x = methods[method]()
10
11       d = {
12           "name": name,
13           "method": method,
14           "nit": x["nit"],
15           "nfev": x["nfev"],
16           "njev": x.get("njev"),
17           "nhev": x.get("nhev"),
18           "success": x["success"]
19           #"message": x["message"]
20       }
21       res.append( pd.DataFrame(d, index=


```


```

1 df = pd.concat([
2     run_collect(name, (1.6, 1.1), cost_f
3     for name, cost_func, arg in zip(
4         ("Well-cond quad", "Ill-cond quad"
5         (mk_quad, mk_quad, mk_rosenbrock),
6         (0.7, 0.02, None)
7     )
8 ])
9
10 df

```

		name	method	nit
nfev	njev	nhev	success	
1	Well-cond quad		CG	2
5	5	None	True	
1	Well-cond quad		newton-cg	5
6	13	0	True	
1	Well-cond quad		newton-cg w/ H	15
15	15	15	True	
1	Well-cond quad		bfgs	8
9	9	None	True	
1	Well-cond quad		bfgs w/o G	8
27	9	None	True	
1	Well-cond quad		l-bfgs-b	6
21	7	None	True	
1	Well-cond quad		nelder-mead	76
147	None	None	True	






Exercise 1

Try minimizing the following function using different optimization methods starting from $x_0 = [0, 0]$, which method(s) appear to work best?

$$f(x) = \exp(x_1 - 1) + \exp(-x_2 + 1) + (x_1 - x_2)^2$$



MVN Example

MVN density cost function

For an n -dimensional multivariate normal we define the $n \times 1$ vectors x and μ and the $n \times n$ covariance matrix Σ ,

$$f(x) = \det(2\pi\Sigma)^{-1/2} \exp\left[-\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu)\right]$$

$$\nabla f(x) = -f(x)\Sigma^{-1} (x - \mu)$$

$$\nabla^2 f(x) = f(x) (\Sigma^{-1} (x - \mu)(x - \mu)^T \Sigma^{-1} - \Sigma^{-1})$$

Our goal will be to find the mode (maximum) of this density.

```
1 def mk_mvn(mu, Sigma):
2     Sigma_inv = np.linalg.inv(Sigma)
3     norm_const = 1 / (np.sqrt(np.linalg.
4
5     # Returns the negative density (since
6     def f(x):
7         x_m = x - mu
8         return -(norm_const *
9                 np.exp(-0.5 * (x_m.T @ Sigma_inv
10
11     def grad(x):
12         return (-f(x) * Sigma_inv @ (x - mu)
13
14     def hess(x):
15         n = len(x)
16         x_m = x - mu
17         return f(x) * ((Sigma_inv @ x_m).r
18
19     return f, grad, hess
```


Gradient checking

One of the most common issues when implementing an optimizer is to get the gradient calculation wrong which can produce problematic results. It is possible to numerically check the gradient function by comparing results between the gradient function and finite differences from the objective function via `optimize.check_grad()`.

```
1 # 2d
2 f, grad, hess = mk_mvn(np.zeros(2), np
3 optimize.check_grad(f, grad, np.zeros(
```

```
np.float64(2.634178031930877e-09)
```

```
1 optimize.check_grad(f, grad, np.ones(2
```

```
np.float64(5.213238144735062e-10)
```

```
1 # 5d
2 f, grad, hess = mk_mvn(np.zeros(5), np
3 optimize.check_grad(f, grad, np.zeros(
```

```
np.float64(2.6031257322754127e-10)
```

```
1 optimize.check_grad(f, grad, np.ones(5
```

```
np.float64(1.725679820308689e-11)
```

```
1 # 10d
2 f, grad, hess = mk_mvn(np.zeros(10), n
3 optimize.check_grad(f, grad, np.zeros(
```

```
np.float64(2.8760747774580336e-12)
```

```
1 optimize.check_grad(f, grad, np.ones(1
```

```
np.float64(2.850398669793798e-14)
```

```
1 # 20d
2 f, grad, hess = mk_mvn(np.zeros(20), n
3 optimize.check_grad(f, grad, np.zeros(
```

```
np.float64(4.965068306494546e-16)
```

```
1 optimize.check_grad(f, grad, np.ones(2
```

```
np.float64(1.0342002372572841e-20)
```

Gradient checking (wrong gradient)

```
1 wrong_grad = lambda x: 2*grad(x)
```

```
1 # 2d
2 f, grad, hess = mk_mvn(np.zeros(2), np
3 optimize.check_grad(f, wrong_grad, [0,
```

np.float64(2.634178031930877e-09)

```
1 optimize.check_grad(f, wrong_grad, [1,
```

np.float64(0.08280196633767578)

```
1 # 5d
2 f, grad, hess = mk_mvn(np.zeros(5), np
3 optimize.check_grad(f, wrong_grad, np.
```

np.float64(2.6031257322754127e-10)

```
1 optimize.check_grad(f, wrong_grad, np.
```

np.float64(0.0018548087267515347)

Hessian checking

Note since the gradient of the gradient / jacobian is the hessian we can use this function to check our implementation of the hessian as well, just use `grad()` as `func` and `hess()` as `grad`.

```
1 # 2d
2 f, grad, hess = mk_mvn(np.zeros(2), np
3 optimize.check_grad(grad, hess, [0,0])
```

`np.float64(3.925231146709438e-17)`

```
1 optimize.check_grad(grad, hess, [1,1])
```

`np.float64(8.399162985270666e-10)`

```
1 # 5d
2 f, grad, hess = mk_mvn(np.zeros(5), np
3 optimize.check_grad(grad, hess, np.zer
```

`np.float64(3.878959614448864e-18)`

```
1 optimize.check_grad(grad, hess, np.one
```

`np.float64(3.8156075963144067e-11)`

Unit MVNs

```

1 df = pd.concat([
2     run_collect(
3         name, np.ones(n), mk_mvn,
4         np.zeros(n), np.eye(n),
5         tol=1e-10,
6         skip=['naive_newton', 'naive_cg',
7     )
8     for name, n in zip(
9         ("5d", "10d", "20d", "30d"),
10        (5, 10, 20, 30)
11    )
12 ])

```


```


1 df

```

	name	method	nit	nfev	njev
	nhev	success			
1	5d	CG	7	52	52
	None	True			
1	5d	newton-cg	3	7	10
0		True			
1	5d	bfgs	3	12	12
	None	True			
1	5d	l-bfgs-b	4	54	9
	None	True			
1	5d	nelder-mead	290	523	None
	None	True			
1	10d	CG	2	24	24
	None	True			
1	10d	newton-cg	2	8	10
2		True			

Performance (Unit MVNs)





Adding correlation


```


1 def build_Sigma(n, corr=0.5):
2     S = np.full((n,n), corr)
3     np.fill_diagonal(S, 1)
4     return S
5
6 df = pd.concat([
7     run_collect(
8         name, np.ones(n), mk_mvn,
9         np.zeros(n), build_Sigma(n),
10        tol=1e-10,
11        skip=['naive_newton', 'naive_cg',
12        ])
13    for name, n in zip(
14        ("5d", "10d", "20d", "30d"),
15        (5, 10, 20, 30)
16    )
17 ])

```

1 df

	name	method	nit	nfev	njev
nhev	success				
1	5d	CG	17	80	80
None	True				
1	5d	newton-cg	4	6	10
0	True				
1	5d	bfgs	4	10	10
None	True				
1	5d	l-bfgs-b	4	54	9
None	True				
1	5d	nelder-mead	224	427	None
None	True				
1	10d	CG	4	31	31
None	True				
1	10d	newton-cg	4	5	9
0	True				





What's going on? (good)

```
1 n = 5
2 f, grad, hess = mk_mvn(np.zeros(n), build_Sigma(n))
```

```
1 optimize.minimize(f, np.ones(n), jac=g
2                   method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
    fun: -0.023337250777292103
       x: [ 1.082e-07  1.059e-07  1.076e-
07  1.088e-07  1.077e-07]
      nit: 14
     jac: [ 8.637e-10  7.556e-10  8.374e-
10  8.917e-10  8.381e-10]
    nfev: 67
    njev: 67
```

```
1 optimize.minimize(f, np.ones(n), jac=g
2                   method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
    fun: -0.023337250777292328
       x: [ 9.586e-10  2.948e-10  7.975e-
10  1.131e-09  8.016e-10]
      nit: 17
     jac: [ 1.376e-11 -1.723e-11  6.238e-
12  2.179e-11  6.430e-12]
    nfev: 80
    njev: 80
```

```
1 optimize.minimize(f, np.ones(n), jac=g
2                   method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
      fun: -0.023337250777292328
         x: [ 6.697e-10  6.565e-10  6.665e-
10  6.731e-10  6.666e-10]
        nit: 18
       jac: [ 5.335e-12  4.720e-12  5.186e-
12  5.494e-12  5.189e-12]
      nfev: 83
      njev: 83
```

What's going on? (okay)

```
1 n = 20
2 f, grad, hess = mk_mvn(np.zeros(n), build_Sigma(n))
```

```
1 optimize.minimize(f, np.ones(n), jac=g
2                   method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
      fun: -2.330191334018497e-06
         x: [-3.221e-04 -3.221e-04 ...
-3.221e-04 -3.221e-04]
        nit: 2
         jac: [-7.148e-11 -7.148e-11 ...
-7.148e-11 -7.148e-11]
        nfev: 27
        njev: 27
```

```
1 optimize.minimize(f, np.ones(n), jac=g
2                   method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
      fun: -2.330191334018497e-06
         x: [-3.221e-04 -3.221e-04 ...
-3.221e-04 -3.221e-04]
        nit: 2
         jac: [-7.148e-11 -7.148e-11 ...
-7.148e-11 -7.148e-11]
        nfev: 27
        njev: 27
```

```
1 optimize.minimize(f, np.ones(n), jac=g
2                   method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
      fun: -2.3301915597315495e-06
         x: [-4.506e-05 -4.506e-05 ...
-4.506e-05 -4.506e-05]
      nit: 2884
      jac: [-9.999e-12 -9.999e-12 ...
-9.999e-12 -9.999e-12]
      nfev: 66313
      njev: 66313
```

What's going on? (bad)

```
1 n = 30
2 f, grad, hess = mk_mvn(np.zeros(n), build_Sigma(n))
```

```
1 optimize.minimize(f, np.ones(n), jac=g
2                     method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
      fun: -2.381146302597316e-09
       x: [ 1.000e+00  1.000e+00 ...
1.000e+00  1.000e+00]
      nit: 0
      jac: [ 1.536e-10  1.536e-10 ...
1.536e-10  1.536e-10]
      nfev: 1
      njev: 1
```

```
1 optimize.minimize(f, np.ones(n), jac=g
2                     method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
      fun: -6.180056227752818e-09
       x: [ 1.203e-01  1.203e-01 ...
1.203e-01  1.203e-01]
      nit: 1
      jac: [ 4.795e-11  4.795e-11 ...
4.795e-11  4.795e-11]
      nfev: 18
      njev: 18
```



```
1 optimize.minimize(f, np.ones(n), jac=g
2                   method="CG", tol=1e-
```

```
message: Optimization terminated
successfully.
success: True
status: 0
      fun: -6.26701117075865e-09
         x: [-5.021e-03 -5.021e-03 ...
-5.021e-03 -5.021e-03]
        nit: 2
         jac: [-2.030e-12 -2.030e-12 ...
-2.030e-12 -2.030e-12]
        nfev: 35
        njev: 35
```

Options (bfgs)

```
1 optimize.show_options(solver="minimize", method="bfgs")
```

Minimization of scalar function of one or more variables using the BFGS algorithm.

Options

disp : bool

Set to True to print convergence messages.

maxiter : int

Maximum number of iterations to perform.

gtol : float

Terminate successfully if gradient norm is less than `gtol`.

norm : float

Order of norm (Inf is max, -Inf is min).

eps : float or ndarray

If `jac is None` the absolute step size used for numerical
approximation of the Jacobian via forward differences

Options (Nelder-Mead)

```
1 optimize.show_options(solver="minimize", method="Nelder-Mead")
```

Minimization of scalar function of one or more variables using the Nelder-Mead algorithm.

Options

disp : bool

Set to True to print convergence messages.

maxiter, maxfev : int

Maximum allowed number of iterations and function evaluations. Will default to ``N*200``, where ``N`` is the number of variables, if neither ``maxiter`` or ``maxfev`` is set. If both ``maxiter`` and ``maxfev`` are set, minimization will stop at the first reached.

return_all : bool, optional

Set to True to return a list of the best solution at each of the iterations.

SciPy implementation

The following code comes from SciPy's `minimize()` implementation:

```
1  if tol is not None:
2      options = dict(options)
3      if meth == 'nelder-mead':
4          options.setdefault('xatol', tol)
5          options.setdefault('fatorl', tol)
6      if meth in ('newton-cg', 'powell', 'tnc'):
7          options.setdefault('xtol', tol)
8      if meth in ('powell', 'l-bfgs-b', 'tnc', 'slsqp'):
9          options.setdefault('ftol', tol)
10     if meth in ('bfgs', 'cg', 'l-bfgs-b', 'tnc', 'dogleg',
11                'trust-ncg', 'trust-exact', 'trust-krylov'):
12         options.setdefault('gtol', tol)
13     if meth in ('cobyla', '_custom'):
14         options.setdefault('tol', tol)
15     if meth == 'trust-constr':
16         options.setdefault('xtol', tol)
17         options.setdefault('gtol', tol)
18         options.setdefault('barrier_tol', tol)
```

Some general advice

- Having access to the gradient is almost always helpful / necessary
- Having access to the hessian can be helpful, but usually does not significantly improve things
- The curse of dimensionality is real
 - Be careful with `tol` - it means different things for different methods
- In general, **BFGS** or **L-BFGS** should be a first choice for most problems (either well- or ill-conditioned)
 - **CG** can perform better for well-conditioned problems with cheap function evaluations

Maximum Likelihood example

Normal MLE

Minimizing

```
1 optimize.minimize(mle_norm, x0=[0,1], method="bfgs")
```

```
message: Desired error not necessarily achieved due to  
precision loss.
```

```
success: False
```

```
status: 2
```

```
fun: nan
```

```
x: [-1.436e+04 -3.533e+03]
```

```
nit: 2
```

```
jac: [      nan      nan]
```

```
hess_inv: [[ 9.443e-01  2.340e-01]  
           [ 2.340e-01  5.905e-02]]
```

```
nfev: 339
```

```
njev: 113
```


Adding constraints

```
1 def mle_norm2(theta):  
2     if theta[1] <= 0:  
3         return np.inf  
4     else:  
5         return -np.sum(norm.logpdf(x, loc=
```

```
1 optimize.minimize(mle_norm2, x0=[0,1],
```

```
message: Optimization terminated  
successfully.  
success: True  
status: 0  
      fun: 163.77575977255518  
       x: [-3.156e+00  1.245e+00]  
      nit: 9  
      jac: [ 0.000e+00  0.000e+00]  
hess_inv: [[ 1.475e-02 -1.179e-04]  
           [-1.179e-04  7.723e-03]]  
      nfev: 43  
      njev: 14
```

Specifying Bounds

It is also possible to specify bounds via `bounds` but this is only available for certain optimization methods.

```
1 optimize.minimize(  
2     mle_norm, x0=[0,1], method="l-bfgs-b",  
3     bounds = [(-1e16, 1e16), (1e-16, 1e16)]  
4 )
```

```
message: CONVERGENCE: RELATIVE REDUCTION OF F <= FACTR*EPSMCH
```

```
success: True
```

```
status: 0
```

```
    fun: 163.77575977287245
```

```
      x: [-3.156e+00  1.245e+00]
```

```
    nit: 10
```

```
    jac: [ 2.046e-04  0.000e+00]
```

```
   nfev: 69
```

```
   njev: 23
```

```
hess_inv: <2x2 LbfgsInvHessProduct with dtype=float64>
```

Exercise 2

Using `optimize.minimize()` recover the shape and scale parameters for these data using MLE.

```
1 from scipy.stats import gamma
2
3 g = gamma(a=2.0, scale=2.0)
4 x = g.rvs(size=100, random_state=1234)
5 x.round(2)
```

```
array([ 4.7 ,  1.11,  1.8 ,  6.19,  3.37,  0.25,  6.45,  0.36,  4.49,
        4.14,  2.84,  1.91,  8.03,  2.26,  2.88,  6.88,  6.84,  6.83,
        6.1 ,  3.03,  3.67,  2.57,  3.53,  2.07,  4.01,  1.51,  5.69,
        3.92,  6.01,  0.82,  2.11,  2.97,  5.02,  9.13,  4.19,  2.82,
       11.81,  1.17,  1.69,  4.67,  1.47, 11.67,  5.25,  3.44,  8.04,
        3.74,  5.73,  6.58,  3.54,  2.4 ,  1.32,  2.04,  2.52,  4.89,
        4.14,  5.02,  4.75,  8.24,  7.6 ,  1.   ,  6.14,  0.58,  2.83,
        2.88,  5.42,  0.5 ,  3.46,  4.46,  1.86,  4.59,  2.24,  2.62,
        3.99,  3.74,  5.27,  1.42,  0.56,  7.54,  5.5 ,  1.58,  5.49,
        6.57,  4.79,  5.84,  8.21,  1.66,  1.53,  4.27,  2.57,  1.48,
        5.23,  3.84,  3.15,  2.1 ,  3.71,  2.79,  0.86,  8.52,  4.36,
        3.3 ])
```