

Optimization (cont.)

Lecture 14

Dr. Colin Rundel

Method Variants

Method: CG in scipy

Scipy's optimize module implements the conjugate gradient algorithm using a variant proposed by Polak and Ribiere, this version does not use the Hessian when calculating the next step. The specific changes are:

- α_k is calculated via a line search along the direction p_k
- and the β_{k+1} calculation is replaced as follows

$$\beta_{k+1} = \frac{r_{k+1}^T \nabla^2 f(x_k) p_k}{p_k^T \nabla^2 f(x_k) p_k} \quad \Rightarrow \quad \beta_{k+1}^{PR} = \frac{\nabla f(x_{k+1}) (\nabla f(x_{k+1}) - \nabla f(x_k))}{\nabla f(x_k)^T \nabla f(x_k)}$$

Method: Newton-CG & BFGS

These are both variants of Newton's method but do not require the Hessian (but can be used by the former if provided).

- Newton-Conjugate Gradient (Newton-CG) algorithm uses a conjugate gradient algorithm to (approximately) invert the local Hessian
- The Broyden-Fletcher-Goldfarb-Shanno (BFGS) algorithm iteratively approximates the inverse Hessian
 - Gradient is also not required and can similarly be approximated using finite differences

Method: Nelder-Mead

This is a gradient free method that uses a series of simplexes which are used to iteratively bracket the minimum.

Method Summary


SciPy Method	Description	Gradient	Hessian
—	Gradient Descent (naive w/ backtracking)	✓	✗
—	Newton's method (naive w/ backtracking)	✓	✓
—	Conjugate Gradient (naive)	✓	✓
"CG"	Nonlinear Conjugate Gradient (Polak and Ribiere variation)	✓	✗
"Newton-CG"	Truncated Newton method (Newton w/ CG step direction)	✓	Optional
"BFGS"	Broyden, Fletcher, Goldfarb, and Shanno (Quasi-newton method)	Optional	✗
"L-BFGS-B"	Limited-memory BFGS (Quasi-newton method)	Optional	✗
"Nelder-Mead"	Nelder-Mead simplex reflection method	✗	✗

Methods collection

```
1 def define_methods(x, f, grad, hess, tol=1e-8):
2     return {
3         "naive_newton": lambda: newtons_method(x, f, grad, hess, tol=tol),
4         "naive_cg": lambda: conjugate_gradient(x, f, grad, hess, tol=tol),
5         "CG": lambda: optimize.minimize(f, x, jac=grad, method="CG", tol=tol),
6         "newton-cg": lambda: optimize.minimize(f, x, jac=grad, hess=None, method="newton-cg", tol=tol),
7         "newton-cg w/ H": lambda: optimize.minimize(f, x, jac=grad, hess=hess, method="newton-cg w/ H", tol=tol),
8         "bfgs": lambda: optimize.minimize(f, x, jac=grad, method="BFGS", tol=tol),
9         "bfgs w/o G": lambda: optimize.minimize(f, x, method="BFGS", tol=tol),
10        "l-bfgs-b": lambda: optimize.minimize(f, x, method="L-BFGS-B", tol=tol),
11        "nelder-mead": lambda: optimize.minimize(f, x, method="Nelder-Mead", tol=tol),
12    }
```


Method Timings

```
1 x = (1.6, 1.1)
2 f, grad, hess = mk_quad(0.7)
3 methods = define_methods(x, f, grad, hess)
4 df = pd.DataFrame({
5     key: timeit.Timer(methods[key]).repeat(10, 100) for key in methods
6 })
```




Timings across cost functions

```
1 def time_cost_func(n, x, name, cost_func, *args):
2     f, grad, hess = cost_func(*args)
3     methods = define_methods(x, f, grad, hess)
4
5     return ( pd.DataFrame({
6         key: timeit.Timer(
7             methods[key]
8         ).repeat(n, n)
9         for key in methods
10    })
11    .melt()
12    .assign(cost_func = name)
13 )
14
15 x = (1.6, 1.1)
16
17 time_cost_df = pd.concat([
18     time_cost_func(10, x, "Well-cond quad", mk_quad, 0.7),
19     time_cost_func(10, x, "Ill-cond quad", mk_quad, 0.02),
20     time_cost_func(10, x, "Rosenbrock", mk_rosenbrock)
21 ])
```



Random starting locations

```
1 pts = np.random.default_rng(seed=1234).uniform(-2,2, (20,2))
2
3 df = pd.concat([
4     pd.concat([
5         time_cost_func(3, x, "Well-cond quad", mk_quad, 0.7),
6         time_cost_func(3, x, "Ill-cond quad", mk_quad, 0.02),
7         time_cost_func(3, x, "Rosenbrock", mk_rosenbrock)
8     ])
9     for x in pts
10 ])
```



Profiling - BFGS (cProfile)

```
1 import cProfile
2
3 f, grad, hess = mk_quad(0.7)
4 def run():
5     for i in range(500):
6         optimize.minimize(fun = f, x0 = (1.6, 1.1), jac=grad, method="BFGS", tol=1e-11)
7
8 cProfile.run('run()', sort="tottime")
```

503504 function calls in 0.209 seconds

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
500	0.043	0.000	0.204	0.000	_optimize.py:1328(_minimize_bfgs)
29000	0.019	0.000	0.019	0.000	{method 'reduce' of 'numpy.ufunc' objects}
13000	0.011	0.000	0.033	0.000	_optimize.py:194(vecnorm)
5000	0.009	0.000	0.021	0.000	<string>:3(f)
18000	0.009	0.000	0.022	0.000	fromnumeric.py:69(_wrapreduction)
5000	0.008	0.000	0.010	0.000	<string>:9(gradient)
4500	0.008	0.000	0.097	0.000	_dcsrc.py:201(__call__)
10000	0.007	0.000	0.018	0.000	numeric.py:2475(array_equal)
4500	0.006	0.000	0.034	0.000	_linesearch.py:86(derphi)
13000	0.005	0.000	0.021	0.000	fromnumeric.py:2338(sum)
5000	0.004	0.000	0.019	0.000	_differentiable_functions.py:39(wrapped)
9000	0.004	0.000	0.005	0.000	_dcsrc.py:269(_iterate)
4500	0.004	0.000	0.104	0.000	_linesearch.py:100(scalar_search_val_for)

Profiling - BFGS (pyinstrument)

```
1 from pyinstrument import Profiler
2
3 f, grad, hess = mk_quad(0.7)
4
5 profiler = Profiler(interval=0.00001)
6
7 profiler.start()
8 opt = optimize.minimize(fun = f, x0 = (1.6, 1.1), jac=grad, method="BFGS", tol=1e-11)
9 p = profiler.stop()
10
11 profiler.write_html("Lec14_bfgs_quad.html")
```


Profiling - Nelder-Mead

```
1 from pyinstrument import Profiler
2
3 f, grad, hess = mk_quad(0.7)
4
5 profiler = Profiler(interval=0.00001)
6
7 profiler.start()
8 opt = optimize.minimize(fun = f, x0 = (1.6, 1.1), method="Nelder-Mead", tol=1e-11)
9 p = profiler.stop()
10
11 profiler.write_html("Lec14_nm_quad.html")
```


optimize.minimize() output

```
1 f, grad, hess = mk_quad(0.7)
```

```
1 optimize.minimize(  
2     fun = f, x0 = (1.6, 1.1),  
3     jac=grad, method="BFGS"  
4 )
```

```
message: Optimization terminated  
successfully.  
success: True  
status: 0  
  fun: 1.2739256453439323e-11  
   x: [-5.318e-07 -8.843e-06]  
  nit: 6  
  jac: [-3.510e-07 -2.860e-06]  
hess_inv: [[ 1.515e+00 -3.438e-03]  
            [-3.438e-03  3.035e+00]]  
  nfev: 7  
  njev: 7
```

```
1 optimize.minimize(  
2     fun = f, x0 = (1.6, 1.1),  
3     jac=grad, hess=hess, method="Newton-CG"  
4 )
```

```
message: Optimization terminated successfully.  
success: True  
status: 0  
  fun: 2.3418652989289317e-12  
   x: [ 0.000e+00  3.806e-06]  
  nit: 11  
  jac: [ 0.000e+00  4.102e-06]  
 nfev: 12  
 njev: 12  
nhev: 11
```

optimize.minimize() output (cont.)

```
1 optimize.minimize(  
2     fun = f, x0 = (1.6, 1.1),  
3     jac=grad, method="CG"  
4 )
```

```
message: Optimization terminated successfully.  
success: True  
status: 0  
  fun: 1.4450021261144105e-32  
   x: [-1.943e-16 -1.110e-16]  
  nit: 2  
  jac: [-1.282e-16 -3.590e-17]  
 nfev: 5  
 njev: 5
```

```
1 optimize.minimize(  
2     fun = f, x0 = (1.6, 1.1),  
3     jac=grad, method="Nelder-Mead"  
4 )
```

```
message: Optimization terminated  
successfully.  
success: True  
status: 0  
  fun: 2.3077013477040082e-10  
   x: [ 1.088e-05  3.443e-05]  
  nit: 46  
 nfev: 89  
final_simplex: (array([[ 1.088e-05,  3.443e-  
05],  
                        [ 1.882e-05, -3.825e-  
05],  
                        [-3.966e-05, -3.147e-  
05]]), array([ 2.308e-10,  3.534e-10,  6.791e-  
10]))
```


Collect

```
1 def run_collect(name, x0, cost_func, *args,
2   f, grad, hess = cost_func(*args)
3   methods = define_methods(x0, f, grad, hess)
4
5   res = []
6   for method in methods:
7       if method in skip: continue
8
9       x = methods[method]()
10      time = timeit.Timer(methods[method]).repeat(10)
11
12      d = {
13          "name": name,
14          "method": method,
15          "nit": x["nit"],
16          "nfev": x["nfev"],
17          "njev": x.get("njev"),
18          "nhev": x.get("nhev"),
19          "success": x["success"],
20          "time": [time]
21      }
22      res.append( pd.DataFrame(d, index=[1]) )
23
24  return pd.concat(res)
```

```
1 df = pd.concat([
2     run_collect(
3         name, (1.6, 1.1),
4         cost_func,
5         arg,
6         skip=['naive_newton', 'naive_cg']
7     )
8     for name, cost_func, arg in zip(
9         ("Well-cond quad", "Ill-cond quad", "Rosenbrock"),
10        (mk_quad, mk_quad, mk_rosenbrock),
11        (0.7, 0.02, None)
12    )
13 ])
```

Runtimes


Function calls



Exercise 1

Try minimizing the following function using different optimization methods starting from $x_0 = [0, 0]$, which method(s) appear to work best?

$$f(x) = \exp(x_1 - 1) + \exp(-x_2 + 1) + (x_1 - x_2)^2$$



MVN Example

MVN density cost function

For an n -dimensional multivariate normal we define the $n \times 1$ vectors x and μ and the $n \times n$ covariance matrix Σ ,

$$f(x) = \det(2\pi\Sigma)^{-1/2} \exp\left[-\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu)\right]$$

$$\nabla f(x) = -f(x)\Sigma^{-1}(x - \mu)$$

$$\nabla^2 f(x) = f(x) (\Sigma^{-1}(x - \mu)(x - \mu)^T \Sigma^{-1} - \Sigma^{-1})$$

Our goal will be to find the mode (maximum) of this density.

```
1 def mk_mvn(mu, Sigma):
2     Sigma_inv = np.linalg.inv(Sigma)
3     norm_const = 1 / (np.sqrt(np.linalg.det(2>
4
5     # Returns the negative density (since we v
6     def f(x):
7         x_m = x - mu
8         return -(norm_const *
9                 np.exp(
10                    -0.5 * x_m.T @ Sigma_inv @ x_m
11                )
12            ).item()
13
14     def grad(x):
15         return (-f(x) * Sigma_inv @ (x - mu))
16
17     def hess(x):
18         n = len(x)
19         x_m = x - mu
20         return f(x) * (
21             (Sigma_inv @ x_m).reshape((n,1))
22             @ (x_m.T @ Sigma_inv).reshape((1,n))
23             - Sigma_inv
24         )
```

Gradient checking

One of the most common issues when implementing an optimizer is to get the gradient calculation wrong which can produce problematic results. It is possible to numerically check the gradient function by comparing results between the gradient function and finite differences from the objective function via `optimize.check_grad()`.

```
1 # 5d
2 f, grad, hess = mk_mvn(
3     np.zeros(5), np.eye(5,5)
4 )
```

```
1 optimize.check_grad(f, grad, np.zeros(5))
```

```
np.float64(2.6031257322754127e-10)
```

```
1 optimize.check_grad(f, grad, np.ones(5))
```

```
np.float64(1.725679820308689e-11)
```

```
1 # 10d
2 f, grad, hess = mk_mvn(
3     np.zeros(10), np.eye(10)
4 )
```

```
1 optimize.check_grad(f, grad, np.zeros(10))
```

```
np.float64(2.8760747774580336e-12)
```

```
1 optimize.check_grad(f, grad, np.ones(10))
```

```
np.float64(2.850398669793798e-14)
```


Gradient checking (wrong gradient)

```
1 wrong_grad = lambda x: 5*grad(x)
```

```
1 # 5d
2 f, grad, hess = mk_mvn(
3     np.zeros(5), np.eye(5)
4 )
```

```
1 optimize.check_grad(f, wrong_grad, np.zeros
```

```
np.float64(2.6031257322754127e-10)
```

```
1 optimize.check_grad(f, wrong_grad, np.ones(!
```

```
np.float64(0.007419234855235744)
```

```
1 # 10d
2 f, grad, hess = mk_mvn(
3     np.zeros(10), np.eye(10)
4 )
```

```
1 optimize.check_grad(f, wrong_grad, np.zeros
```

```
np.float64(2.8760747774580336e-12)
```

```
1 optimize.check_grad(f, wrong_grad, np.ones(!
```

```
np.float64(8.703385925704049e-06)
```

Why does `np.ones()` detect an issue but `np.zeros()` does not?

Hessian checking

Note since the gradient of the gradient is the hessian we can use this function to check our implementation of the hessian as well, just use `grad()` as `func` and `hess()` as `grad`.

```
1 # 5d
2 f, grad, hess = mk_mvn(np.zeros(5), np.eye(5))
```

```
1 optimize.check_grad(grad, hess, np.zeros(5))
```

```
np.float64(3.878959614448864e-18)
```

```
1 optimize.check_grad(grad, hess, np.ones(5))
```

```
np.float64(3.8156075963144067e-11)
```

```
1 # 10d
2 f, grad, hess = mk_mvn(np.zeros(10), np.eye(10))
```

```
1 optimize.check_grad(grad, hess, np.zeros(10))
```

```
np.float64(4.2856853864461685e-20)
```

```
1 optimize.check_grad(grad, hess, np.ones(10))
```

```
np.float64(8.551196009381392e-14)
```


Unit MVNs

```
1 rng = np.random.default_rng(seed=1234)
2 runif = rng.uniform(-1,1, size=25)
3
4 df = pd.concat([
5     run_collect(
6         name, runif[:n], mk_mvn,
7         np.zeros(n), np.eye(n),
8         tol=1e-10,
9         skip=['naive_newton', 'naive_cg', 'bfgs w/o G', 'newton-cg w/ H', 'l-bfgs-b', 'nelder-mea
10     ])
11     for name, n in zip(
12         ("5d", "10d", "25d"),
13         (5, 10, 25)
14     )
15 ])
```

Performance (Unit MVNs)

Run times

Function calls



What's going on? (good)

```
1 n = 5
2 f, grad, hess = mk_mvn(np.zeros(n), np.eye(n))
```

```
1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="newton-cg", tol=1e-10
4 )
```

```
message: Optimization terminated successfully.
success: True
status: 0
  fun: -0.010105326013811651
    x: [ 5.071e-11 -1.274e-11  4.502e-11
-2.535e-11 -1.924e-11]
   nit: 4
   jac: [ 5.125e-13 -1.288e-13  4.550e-13
-2.562e-13 -1.945e-13]
  nfev: 5
  njev: 9
  nhev: 0
```

```
1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="bfgs", tol=1e-10
4 )
```

```
message: Optimization terminated
successfully.
success: True
status: 0
  fun: -0.010105326013811651
    x: [-2.482e-13  6.237e-14 -2.203e-13
1.240e-13  9.422e-14]
   nit: 4
   jac: [-2.508e-15  6.303e-16 -2.227e-15
1.253e-15  9.521e-16]
  hess_inv: [[ 4.463e+01 -1.096e+01 ...
-2.181e+01 -1.656e+01]
  [-1.096e+01  3.756e+00 ...
5.481e+00  4.161e+00]
  ...
  [-2.181e+01  5.481e+00 ...
1.190e+01  8.276e+00]
  [-1.656e+01  4.161e+00 ...
```

What's going on? (okay)

```
1 n = 10
2 f, grad, hess = mk_mvn(np.zeros(n), np.eye(n))
```

```
1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="newton-cg", tol=1e-10
4 )
```

```
message: Optimization terminated successfully.
success: True
status: 0
  fun: -0.00010211745783654051
    x: [-8.223e-04  2.067e-04 -7.301e-04
 4.111e-04  3.120e-04
        6.588e-04  4.454e-04  3.130e-04
-8.005e-04  4.077e-04]
  nit: 3
   jac: [-8.397e-08  2.110e-08 -7.455e-08
 4.198e-08  3.186e-08
        6.727e-08  4.549e-08  3.196e-08
-8.174e-08  4.163e-08]
 nfev: 6
 njev: 9
nhev: 0
```

```
1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="bfgs", tol=1e-10
4 )
```

```
message: Optimization terminated
successfully.
success: True
status: 0
  fun: -0.00010211761384541865
    x: [-2.347e-09  5.898e-10 -2.084e-09
 1.173e-09  8.906e-10
        1.880e-09  1.271e-09  8.933e-10
-2.285e-09  1.164e-09]
  nit: 2
   jac: [-2.396e-13  6.023e-14 -2.128e-13
 1.198e-13  9.094e-14
        1.920e-13  1.298e-13  9.122e-14
-2.333e-13  1.188e-13]
 hess_inv: [[ 1.685e+04 -4.235e+03 ...
 1.641e+04 -8.356e+03]
           [-4.235e+03  1.065e+03 ...
-4.122e+03  2.100e+03]
```

What's going on? (bad)

```
1 n = 25
2 f, grad, hess = mk_mvn(np.zeros(n), np.eye(n))
```

```
1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="newton-cg", tol=1e-10
4 )
```

```
message: Optimization terminated successfully.
success: True
status: 0
fun: -1.2867324357023428e-12
x: [ 9.534e-01 -2.396e-01 ... 2.220e-01
-8.797e-01]
nit: 1
jac: [ 1.227e-12 -3.083e-13 ... 2.857e-13
-1.132e-12]
nfev: 1
njev: 1
nhev: 0
```

```
1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="bfgs", tol=1e-10
4 )
```

```
message: Optimization terminated
successfully.
success: True
status: 0
fun: -1.2867324357023428e-12
x: [ 9.534e-01 -2.396e-01 ... 2.220e-
01 -8.797e-01]
nit: 0
jac: [ 1.227e-12 -3.083e-13 ... 2.857e-
13 -1.132e-12]
hess_inv: [[1 0 ... 0 0]
[0 1 ... 0 0]
...
[0 0 ... 1 0]
[0 0 ... 0 1]]
nfev: 1
njev: 1
```

All bad?

```
1 optimize.minimize(  
2     f, runif[:n], jac=grad,  
3     method="nelder-mead", tol=1e-10  
4 )
```

```
message: Maximum number of function evaluations has been exceeded.  
success: False  
status: 1  
  fun: -5.2161537392613975e-11  
   x: [ 7.181e-02 -3.136e-01 ...  3.193e-01  3.222e-02]  
  nit: 4136  
 nfev: 5000  
final_simplex: (array([[ 7.181e-02, -3.136e-01, ...,  3.193e-01,  
                        3.222e-02],  
                      [ 7.275e-02, -3.237e-01, ...,  3.218e-01,  
                        2.192e-02],  
                      ...,  
                      [ 8.238e-02, -3.143e-01, ...,  3.247e-01,  
                        2.232e-02],  
                      [ 8.105e-02, -3.178e-01, ...,  3.119e-01,  
                        3.078e-02]], shape=(26, 25)), array([-5.216e-11, -5.216e-11, ...,  
-5.213e-11, -5.213e-11],  
                  shape=(26,)))
```


Options (newton-cg)

```
1 optimize.show_options(solver="minimize", method="newton-cg")
```

Minimization of scalar function of one or more variables using the Newton-CG algorithm.

Note that the ``jac`` parameter (Jacobian) is required.

Options

`disp` : bool

Set to True to print convergence messages.

`xtol` : float

Average relative error in solution ``xopt`` acceptable for convergence.

`maxiter` : int

Maximum number of iterations to perform.

`eps` : float or ndarray

If ``hessp`` is approximated, use this value for the step size.

`return_all` : bool, optional

Set to True to return a list of the best solution at each of the

Options (bfgs)

```
1 optimize.show_options(solver="minimize", method="bfgs")
```

Minimization of scalar function of one or more variables using the BFGS algorithm.

Options

disp : bool

Set to True to print convergence messages.

maxiter : int

Maximum number of iterations to perform.

gtol : float

Terminate successfully if gradient norm is less than `gtol`.

norm : float

Order of norm (Inf is max, -Inf is min).

eps : float or ndarray

If `jac is None` the absolute step size used for numerical approximation of the jacobian via forward differences.

return_all : bool, optional

Set to True to return a list of the best solution at each of the

Options (Nelder-Mead)

```
1 optimize.show_options(solver="minimize", method="newton-cg")
```

Minimization of scalar function of one or more variables using the Newton-CG algorithm.

Note that the `jac` parameter (Jacobian) is required.

Options

disp : bool

Set to True to print convergence messages.

xtol : float

Average relative error in solution `xopt` acceptable for convergence.

maxiter : int

Maximum number of iterations to perform.

eps : float or ndarray

If `hessp` is approximated, use this value for the step size.

return_all : bool, optional

Set to True to return a list of the best solution at each of the

SciPy implementation

The following code comes from SciPy's `minimize()` implementation:

```
1  if tol is not None:
2      options = dict(options)
3      if meth == 'nelder-mead':
4          options.setdefault('xatol', tol)
5          options.setdefault('fatol', tol)
6      if meth in ('newton-cg', 'powell', 'tnc'):
7          options.setdefault('xtol', tol)
8      if meth in ('powell', 'l-bfgs-b', 'tnc', 'slsqp'):
9          options.setdefault('ftol', tol)
10     if meth in ('bfgs', 'cg', 'l-bfgs-b', 'tnc', 'dogleg',
11                'trust-ncg', 'trust-exact', 'trust-krylov'):
12         options.setdefault('gtol', tol)
13     if meth in ('cobyla', '_custom'):
14         options.setdefault('tol', tol)
15     if meth == 'trust-constr':
16         options.setdefault('xtol', tol)
17         options.setdefault('gtol', tol)
18         options.setdefault('barrier_tol', tol)
```

Fixing our tolerances?

```
1 n = 25
2 f, grad, hess = mk_mvn(np.zeros(n), np.eye(n))
```

```
1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="newton-cg", tol=1e-16
4 )
```

```
message: Optimization terminated successfully.
success: True
status: 0
  fun: -1.2867324357023428e-12
    x: [ 9.534e-01 -2.396e-01 ...  2.220e-01
-8.797e-01]
   nit: 1
  jac: [ 1.227e-12 -3.083e-13 ...  2.857e-13
-1.132e-12]
 nfev: 1
njev: 1
nhev: 0
```

```
1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="bfgs", tol=1e-16
4 )
```

```
message: Optimization terminated
successfully.
success: True
status: 0
  fun: -1.0537841099018322e-10
    x: [-2.645e-08  6.648e-09 ... -6.160e-
09  2.441e-08]
   nit: 3
  jac: [-2.788e-18  7.006e-19 ... -6.492e-
19  2.572e-18]
hess_inv: [[ 9.790e+08 -2.461e+08 ...
2.280e+08 -9.034e+08]
          [-2.461e+08  6.184e+07 ...
-5.730e+07  2.270e+08]
          ...
          [ 2.280e+08 -5.730e+07 ...
5.310e+07 -2.104e+08]
          [ 0.034e+08  2.270e+08
```

Limits of floating point precision

Every type of floating point value has finite precision due to the limitations of how it is represented. This value is typically referred to as the machine epsilon value, this is the smallest possible spacing between 1.0 and the next representable floating-point number.

```
1 np.finfo(np.float64).eps
```

```
np.float64(2.220446049250313e-16)
```

```
1 np.finfo(np.float32).eps
```

```
np.float32(1.1920929e-07)
```

```
1 np.finfo(np.float16).eps
```

```
np.float16(0.000977)
```

```
1 1+np.finfo(np.float64).eps > 1
```

```
np.True_
```

```
1 1+np.finfo(np.float64).eps/2 > 1
```

```
np.False_
```

Fixes?

```
1 def mk_prop_mvn(mu, Sigma):
2     Sigma_inv = np.linalg.inv(Sigma)
3     #norm_const = 1 / (np.sqrt(np.linalg.det(2*np.pi*Sigma)))
4     norm_const = 1
5
6     # Returns the negative density (since we want the max not min)
7     def f(x):
8         x_m = x - mu
9         return -(norm_const *
10                 np.exp(
11                     -0.5 * x_m.T @ Sigma_inv @ x_m
12                 )
13                 ).item()
14
15     def grad(x):
16         return (-f(x) * Sigma_inv @ (x - mu))
17
18     def hess(x):
19         n = len(x)
20         x_m = x - mu
21         return f(x) * (
22             (Sigma_inv @ x_m).reshape((n,1))
23             @ (x_m.T @ Sigma_inv).reshape((1,n))
24             - Sigma_inv
```

```

1 n = 25
2 f, grad, hess = mk_prop_mvn(np.zeros(n), np.eye(n))

```

```

1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="newton-cg", tol=1e-10
4 )

```

```

message: Optimization terminated successfully.
success: True
status: 0
      fun: -1.0
         x: [-3.701e-16  9.302e-17 ... -8.619e-17
3.415e-16]
        nit: 4
       jac: [-3.701e-16  9.302e-17 ... -8.619e-17
3.415e-16]
      nfev: 10
      njev: 14
      nhev: 0

```

```

1 optimize.minimize(
2     f, runif[:n], jac=grad,
3     method="bfgs", tol=1e-10
4 )

```

```


message: Optimization terminated
successfully.
success: True
status: 0
      fun: -1.0
         x: [-3.040e-15  7.639e-16 ... -7.079e-
16  2.805e-15]
        nit: 4
       jac: [-3.040e-15  7.639e-16 ... -7.079e-
16  2.805e-15]
      hess_inv: [[ 1.000e+00 -6.660e-09 ...  6.171e-
09 -2.445e-08]
                [-6.660e-09  1.000e+00 ... -1.551e-
09  6.145e-09]
                ...
                [ 6.171e-09 -1.551e-09 ...
1.000e+00 -5.694e-09]
                [ 2.445e-08  6.145e-09 ... 5.694e-

```


Performance

Run times

Function calls



Some general advice

- Having access to the gradient is almost always helpful / necessary
- Having access to the hessian can be helpful, but usually does not significantly improve things
- The curse of dimensionality is real
 - Be careful with `tol` - it means different things for different methods
- In general, **BFGS** or **L-BFGS** should be a first choice for most problems (either well- or ill-conditioned)
 - **CG** can perform better for well-conditioned problems with cheap function evaluations

Maximum Likelihood example

Normal MLE

```
1 from scipy.stats import norm
2
3 n = norm(-3.2, 1.25)
4 x = n.rvs(size=100, random_state=1234)
5 {'μ': x.mean(), 'σ': x.std()}
```

```
{'μ': np.float64(-3.156109646093205), 'σ':
np.float64(1.2446060629192537)}
```

```
1 mle_norm = lambda θ: -np.sum(
2     norm.logpdf(x, loc=θ[0], scale=θ[1])
3 )
```

```
1 mle_norm([0,1])
```

```
np.float64(667.3974708213642)
```

```
1 mle_norm([-3, 1])
```

```
np.float64(170.56457699340282)
```

```
1 mle_norm([-3.2, 1.25])
```

```
np.float64(163.83926813257395)
```

```
1 mle_norm([-3.3, 1.25])
```

```
np.float64(164.44016639757749)
```

Minimizing

```
1 optimize.minimize(mle_norm, x0=[0,1], method="bfgs")
```

```
message: Desired error not necessarily achieved due to precision loss.
success: False
status: 2
  fun: nan
    x: [-1.436e+04 -3.533e+03]
  nit: 2
  jac: [      nan      nan]
hess_inv: [[ 9.443e-01  2.340e-01]
           [ 2.340e-01  5.905e-02]]
  nfev: 339
  njev: 113
```

Adding constraints

```
1 def mle_norm2(theta):
2     if theta[1] <= 0:
3         return np.inf
4     else:
5         return -np.sum(
6             norm.logpdf(x, loc=theta[0], scale=theta[1])
7         )
```

```
1 optimize.minimize(mle_norm2, x0=[0,1], method='Nelder-Mead')
```

```
message: Optimization terminated
successfully.
success: True
status: 0
fun: 163.77575977255518
x: [-3.156e+00  1.245e+00]
nit: 9
jac: [ 0.000e+00  0.000e+00]
hess_inv: [[ 1.475e-02 -1.179e-04]
            [-1.179e-04  7.723e-03]]
nfev: 43
njev: 14
```

```
1 {'μ': x.mean(), 'σ': x.std()}
```

```
{'μ': np.float64(-3.156109646093205), 'σ': np.float64(1.2446060629192537)}
```

Specifying Bounds

It is also possible to specify bounds via `bounds` but this is only available for certain optimization methods (i.e. Nelder-Mead & L-BFGS-B).

```
1 optimize.minimize(  
2     mle_norm, x0=[0,1], method="l-bfgs-b",  
3     bounds = [(-1e16, 1e16), (1e-16, 1e16)]  
4 )
```

```
message: CONVERGENCE: RELATIVE REDUCTION OF F <= FACTR*EPSMCH
```

```
success: True
```

```
status: 0
```

```
    fun: 163.77575977287245
```

```
      x: [-3.156e+00  1.245e+00]
```

```
    nit: 10
```

```
    jac: [ 2.046e-04  0.000e+00]
```

```
   nfev: 69
```

```
   njev: 23
```

```
hess_inv: <2x2 LbfgsInvHessProduct with dtype=float64>
```

Exercise 2

Using `optimize.minimize()` recover the shape and scale parameters for these data using MLE.

```
1 from scipy.stats import gamma
2
3 g = gamma(a=2.0, scale=2.0)
4 x = g.rvs(size=100, random_state=1234)
5 x.round(2)
```

```
array([ 4.7 ,  1.11,  1.8 ,  6.19,  3.37,  0.25,  6.45,  0.36,  4.49,
        4.14,  2.84,  1.91,  8.03,  2.26,  2.88,  6.88,  6.84,  6.83,
        6.1 ,  3.03,  3.67,  2.57,  3.53,  2.07,  4.01,  1.51,  5.69,
        3.92,  6.01,  0.82,  2.11,  2.97,  5.02,  9.13,  4.19,  2.82,
       11.81,  1.17,  1.69,  4.67,  1.47, 11.67,  5.25,  3.44,  8.04,
        3.74,  5.73,  6.58,  3.54,  2.4 ,  1.32,  2.04,  2.52,  4.89,
        4.14,  5.02,  4.75,  8.24,  7.6 ,  1.   ,  6.14,  0.58,  2.83,
        2.88,  5.42,  0.5 ,  3.46,  4.46,  1.86,  4.59,  2.24,  2.62,
        3.99,  3.74,  5.27,  1.42,  0.56,  7.54,  5.5 ,  1.58,  5.49,
        6.57,  4.79,  5.84,  8.21,  1.66,  1.53,  4.27,  2.57,  1.48,
        5.23,  3.84,  3.15,  2.1 ,  3.71,  2.79,  0.86,  8.52,  4.36,
        3.3 ])
```